



# ITH508 컴퓨터망

**Transport Layer**

**Hwangnam Kim**  
**hnkim@korea.ac.kr**  
**School of Electrical Engineering**  
**Korea University**

# Transport Layer



## ■ Principles behind transport layer services:

- ▶ Multiplexing/demultiplexing
- ▶ Reliable data transfer
- ▶ Flow control
- ▶ Congestion control

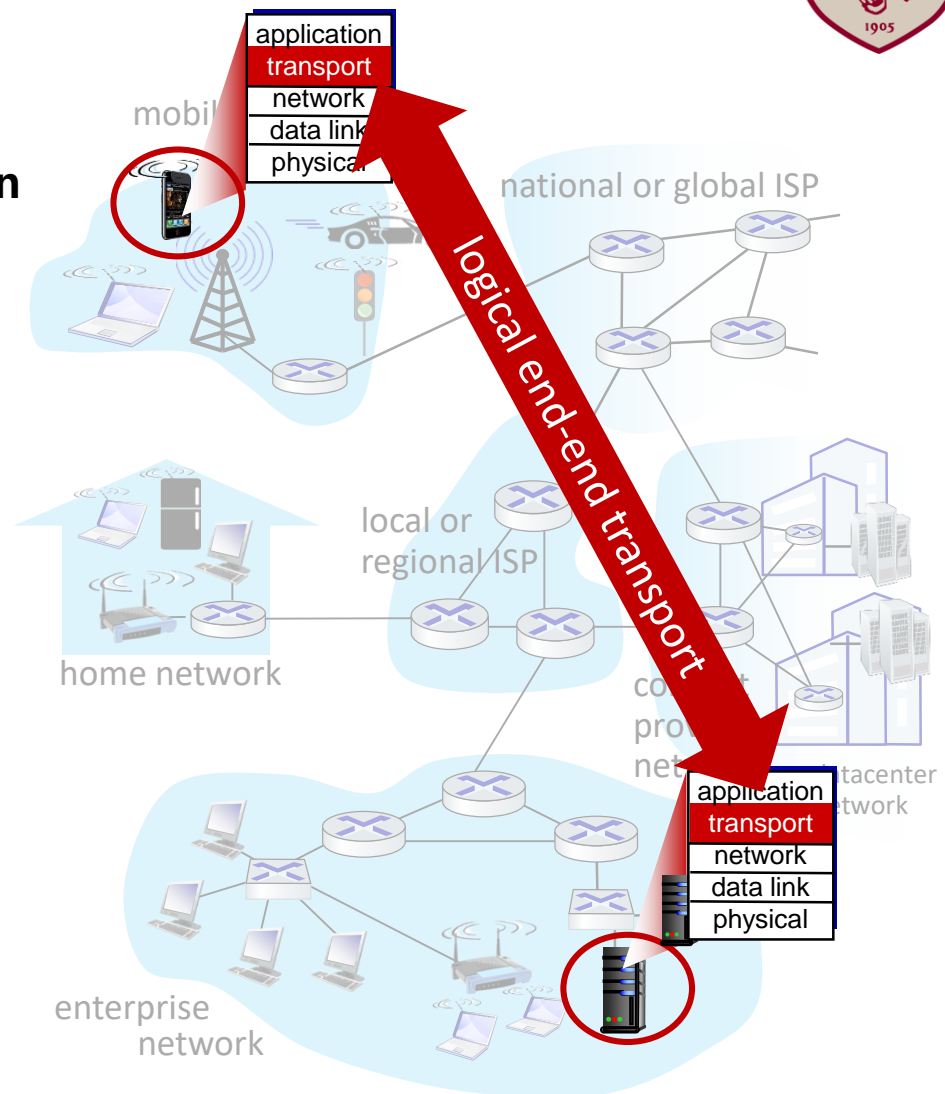
## ■ Transport layer protocols in the Internet:

- ▶ UDP: connectionless transport
- ▶ TCP: connection-oriented transport
- ▶ TCP congestion control

# Transport Services and Protocols



- Provide **logical communication** between **app processes** running on different hosts
- Transport protocols run in end systems
  - ▶ Sender side:
    - Breaks app messages into **segments**
    - Passes to network layer
  - ▶ Receiver side:
    - Reassembles **segments** into messages
    - Passes to app layer
- More than one transport protocol available to apps
  - ▶ Internet: TCP and UDP



# Transport vs. Network Layer



- *Network layer:* logical communication **between hosts**
- *Transport layer:* logical communication **between processes**
  - ▶ Relies on, enhances, network layer services

# Internet Transport-layer Protocols



## ■ Reliable, in-order delivery: TCP

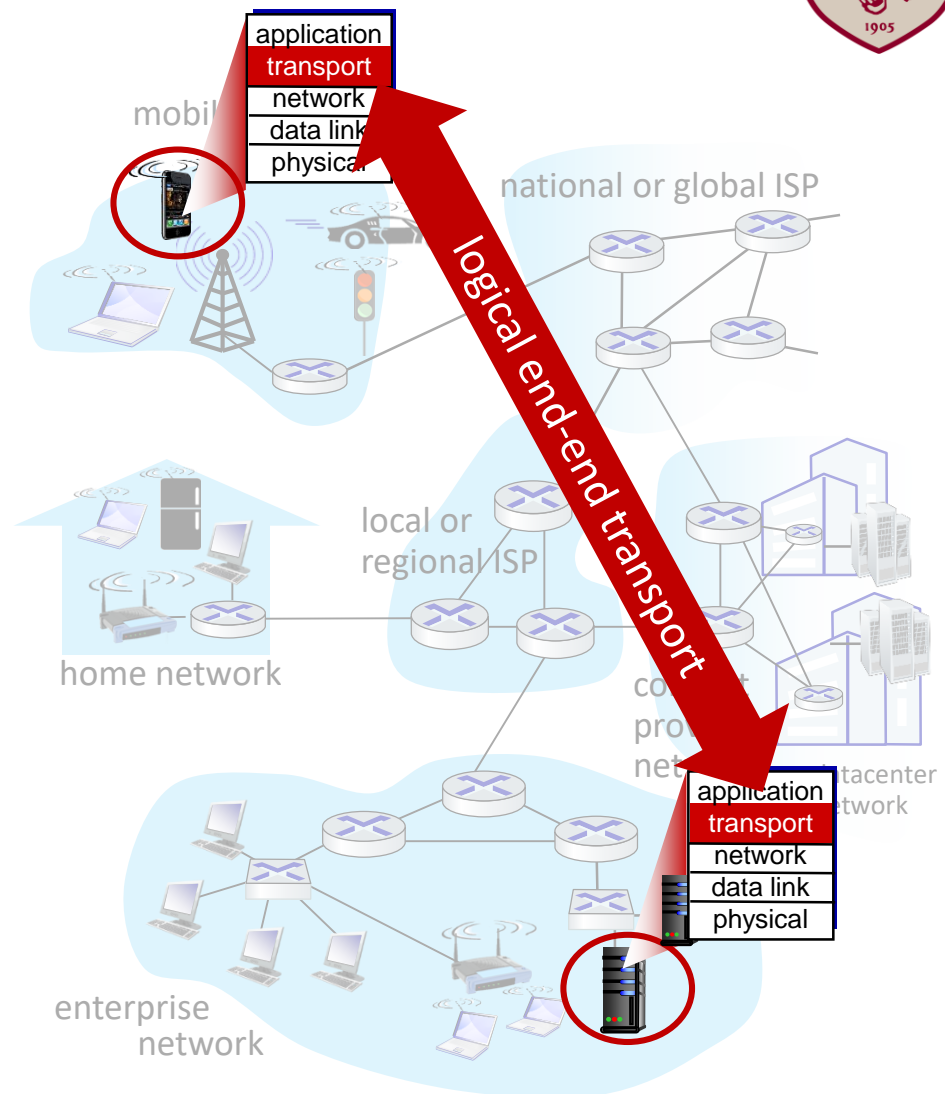
- ▶ Congestion control
- ▶ Flow control
- ▶ Error control
- ▶ Connection setup

## ■ Unreliable, unordered delivery: UDP

- ▶ Simple extension of "best-effort" IP

## ■ Services not available:

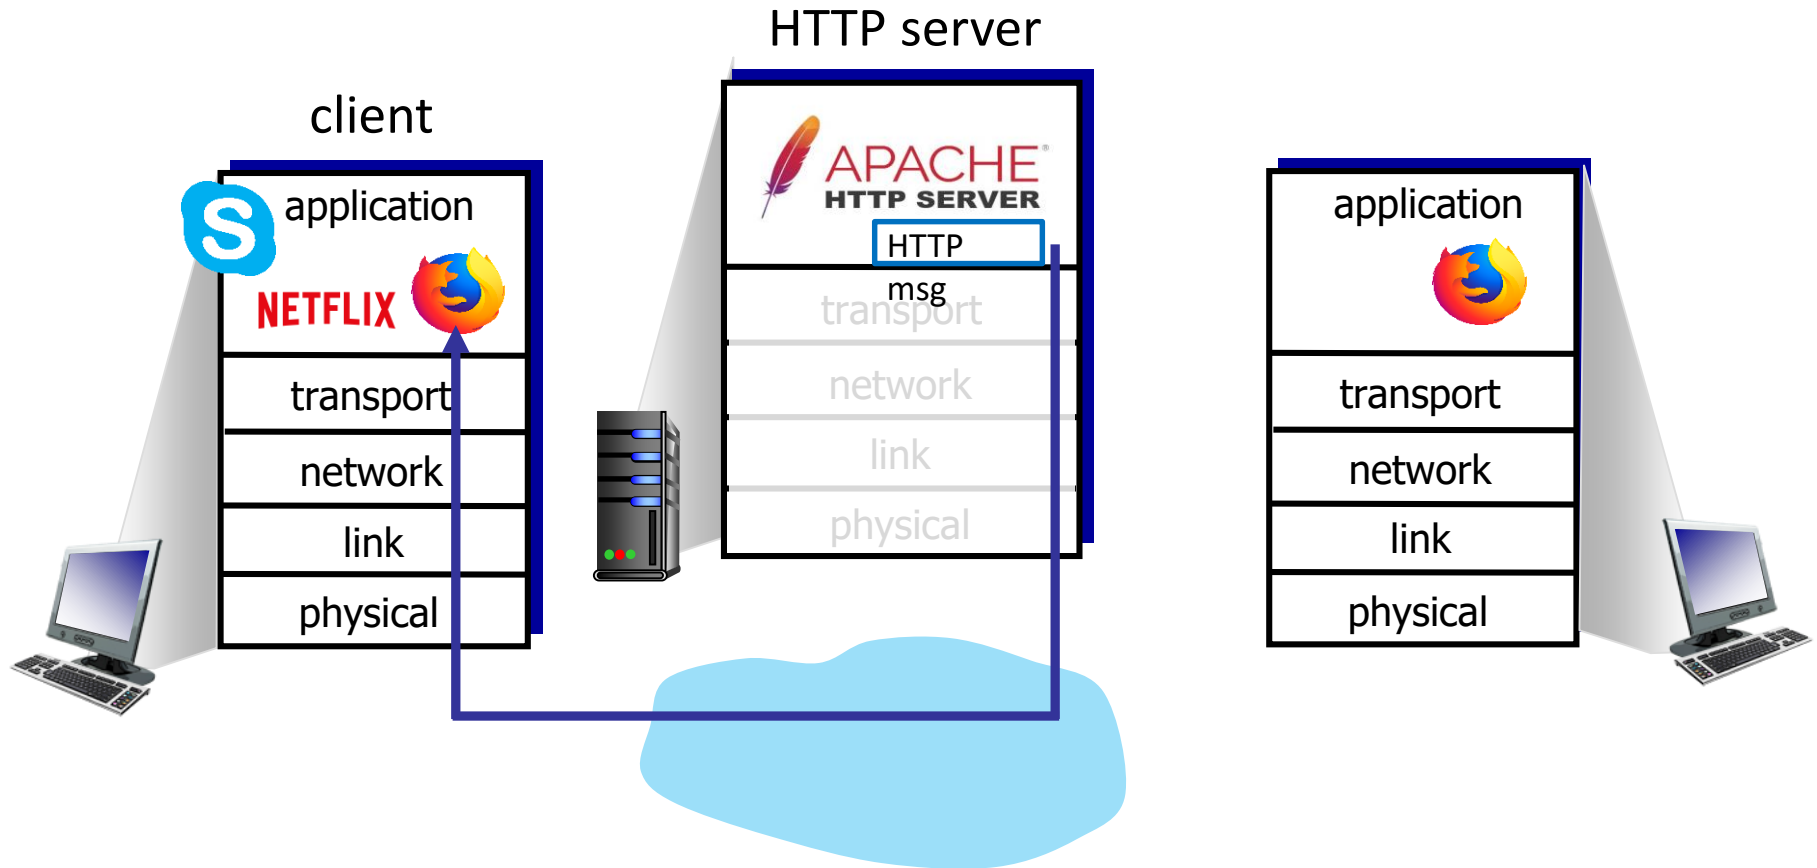
- ▶ Delay guarantees
- ▶ Bandwidth guarantees



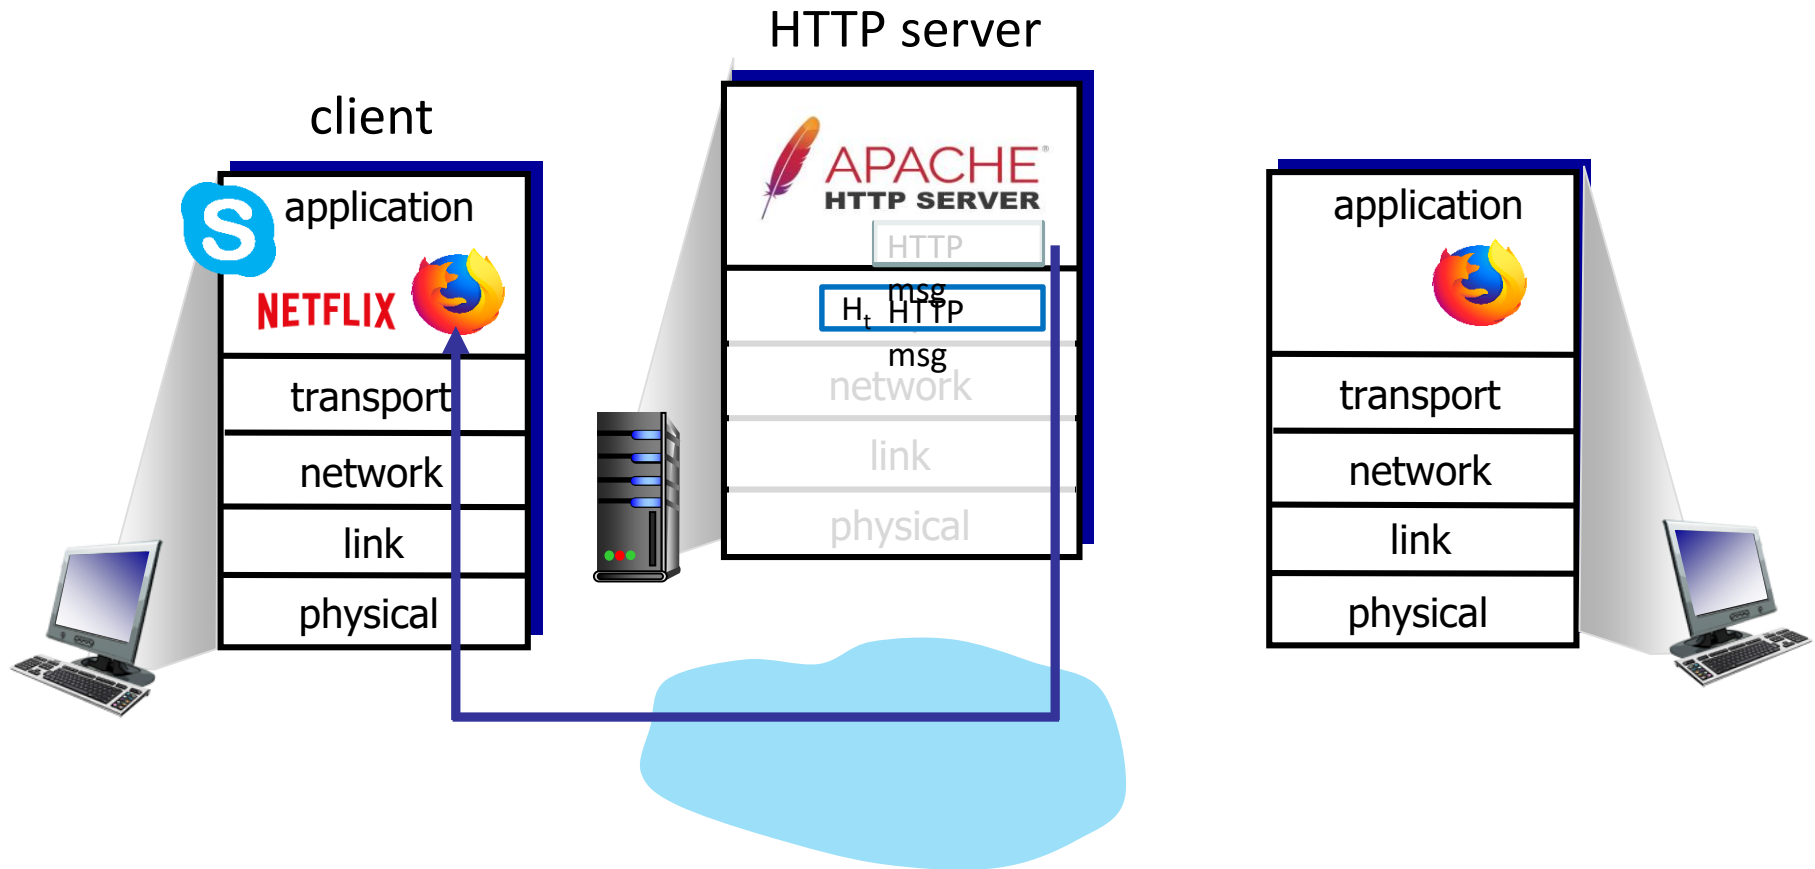
Common feature

# MULTIPLEXING/DEMULTIPLEXI NG

# Multiplexing/Demultiplexing

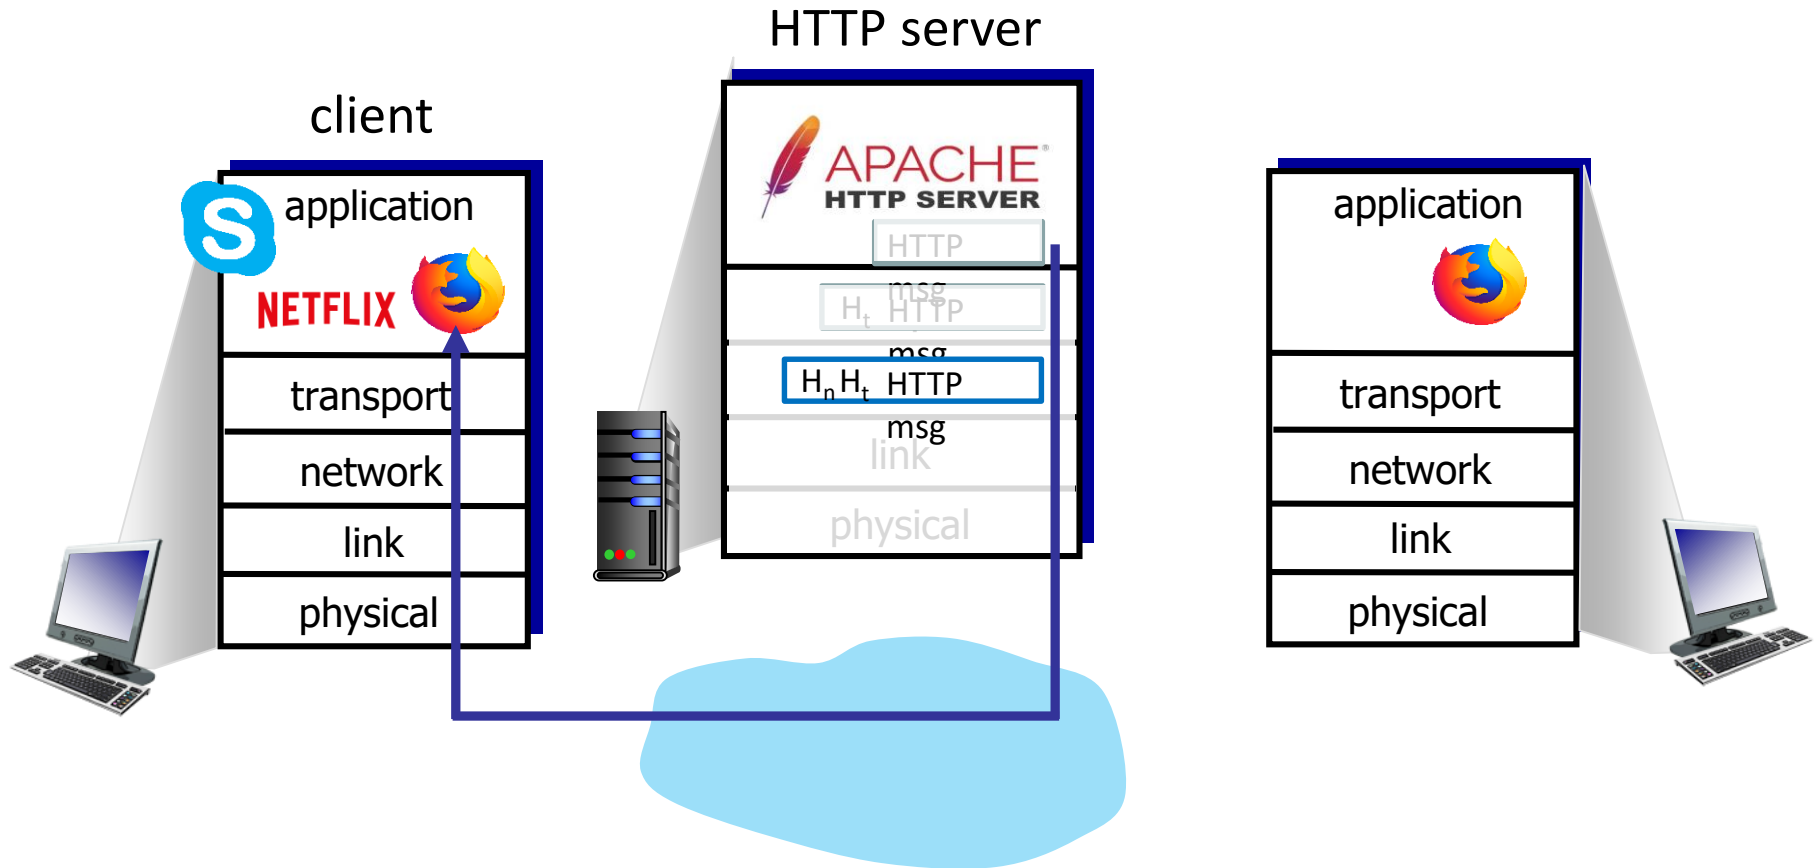


# Multiplexing/Demultiplexing

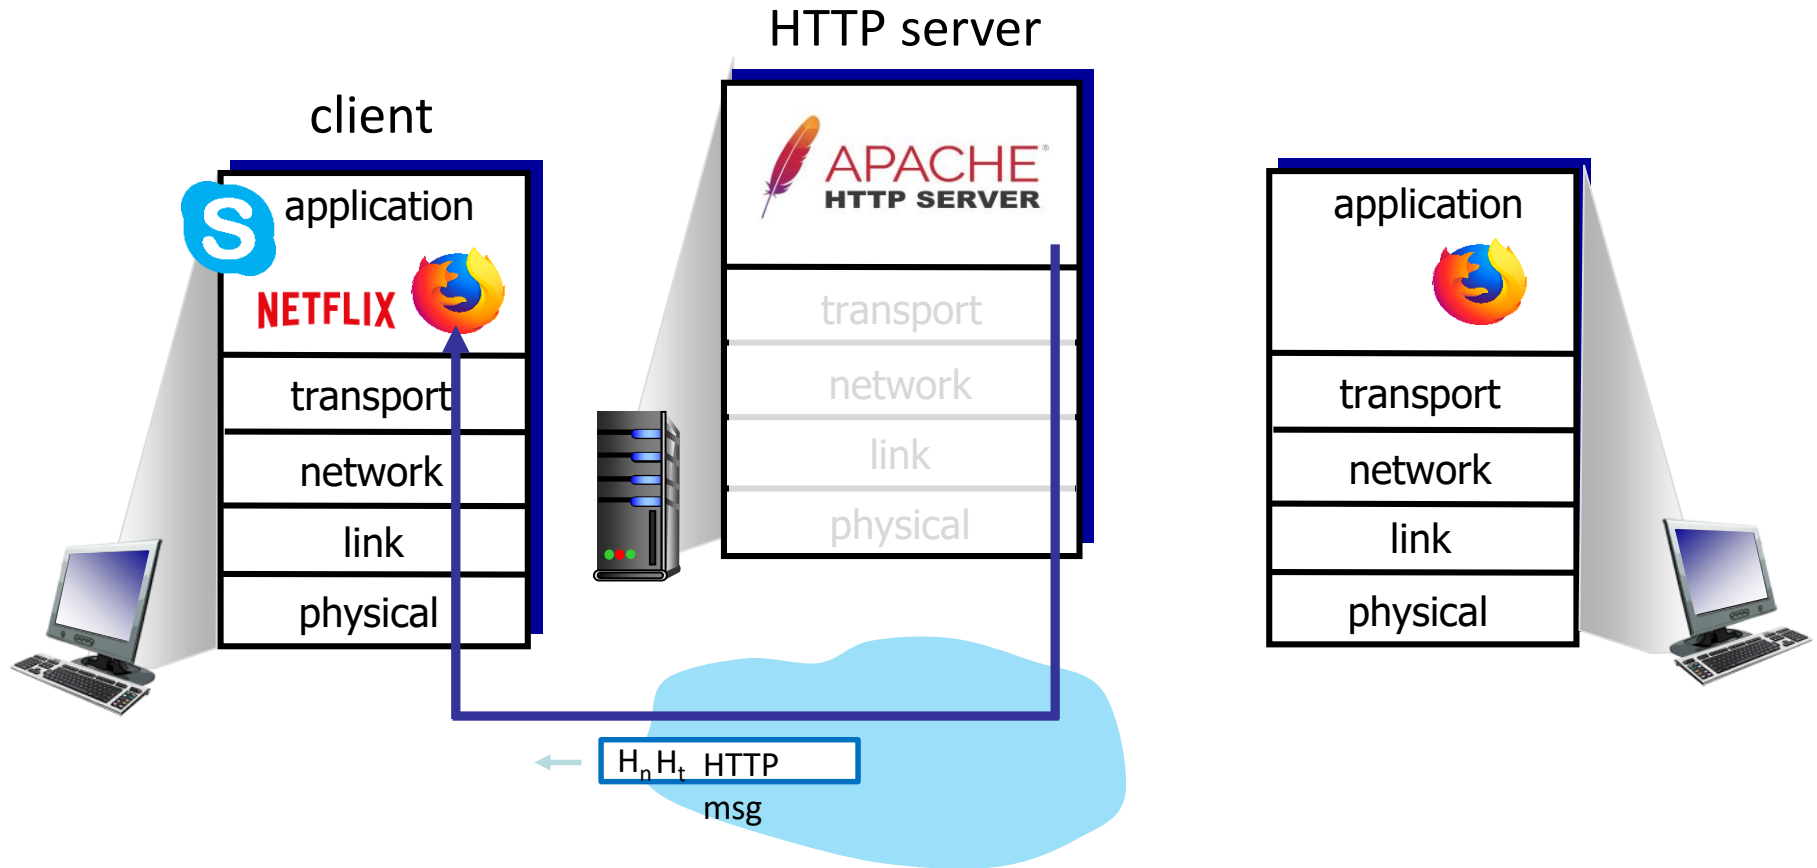




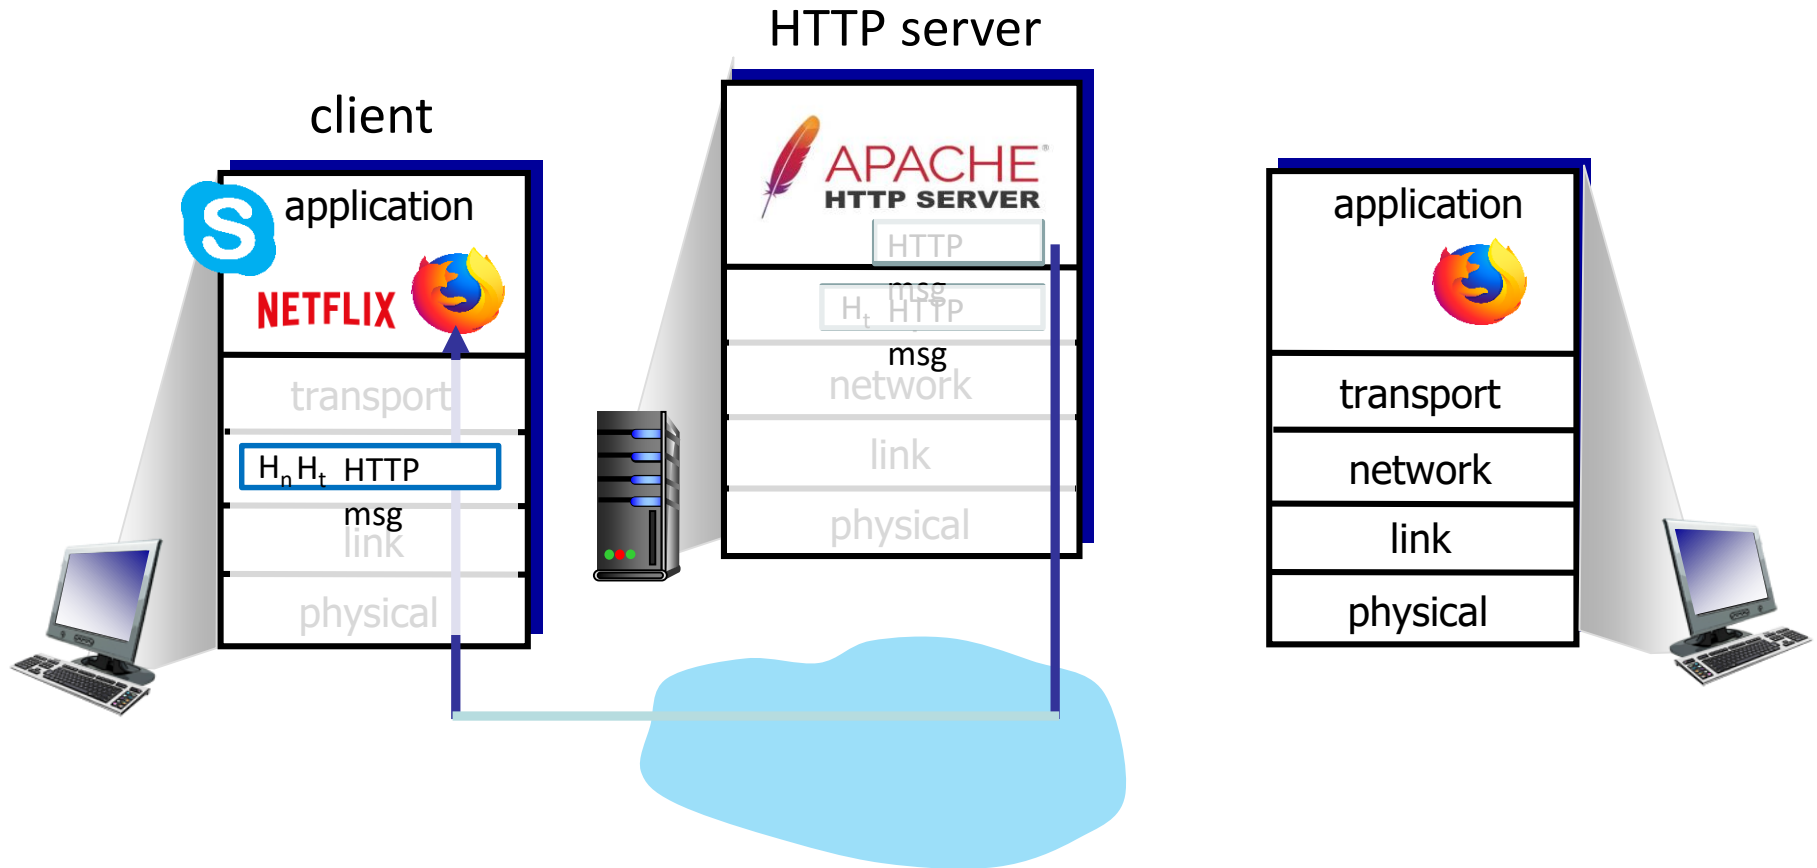
# Multiplexing/Demultiplexing



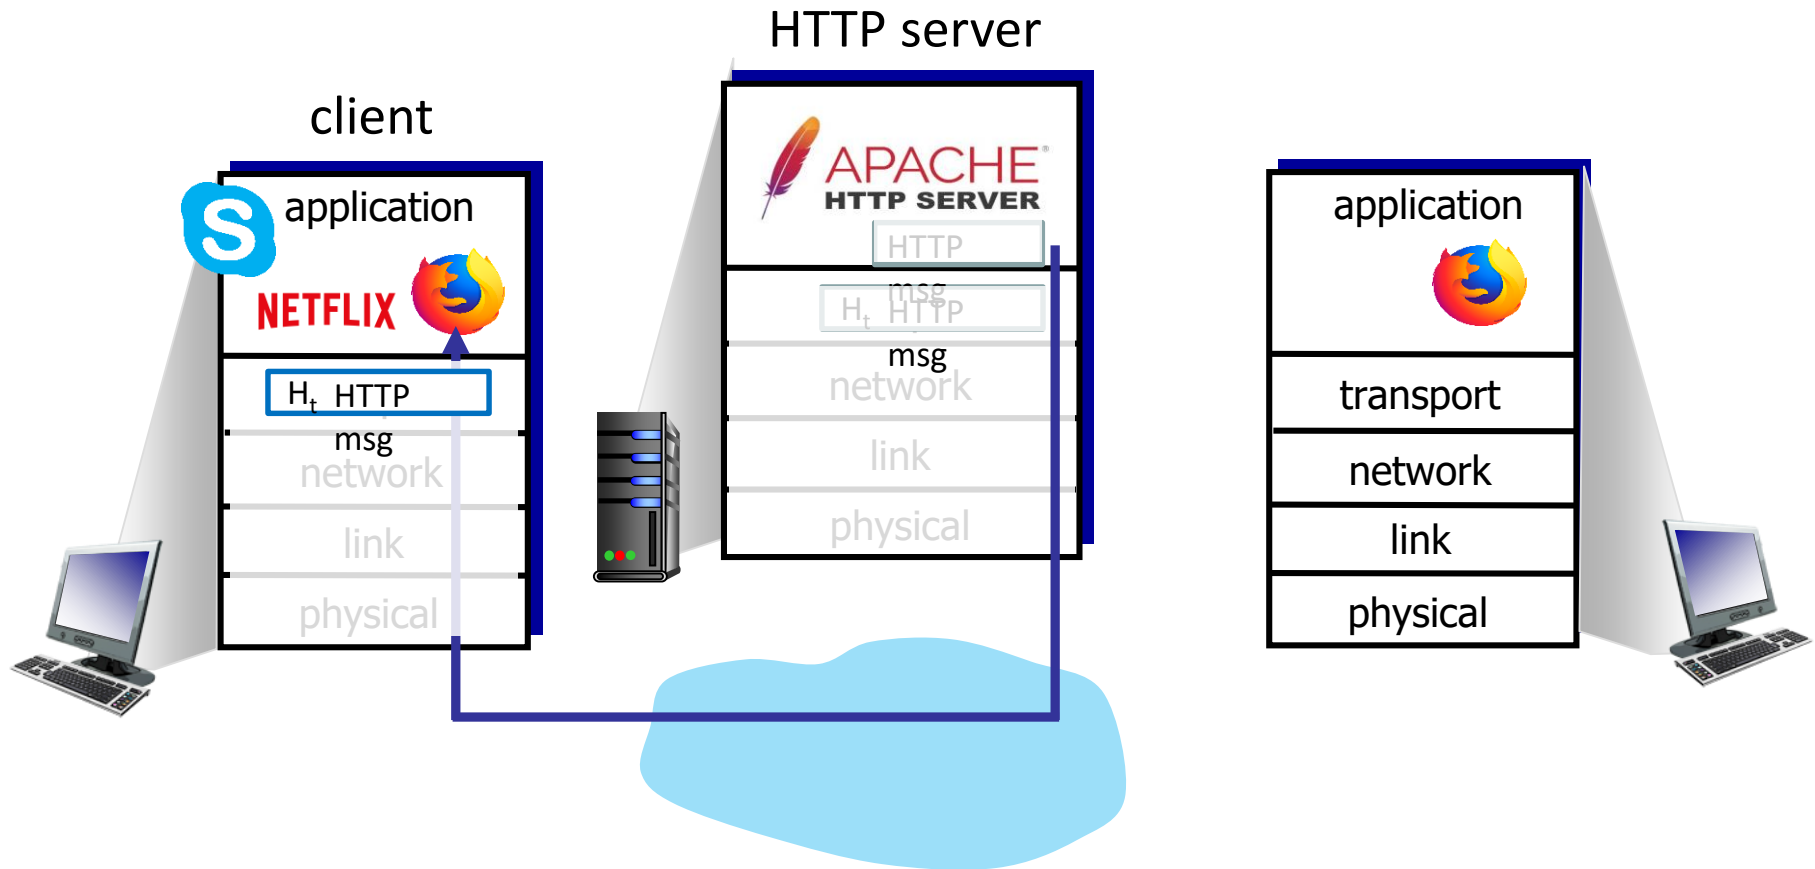
# Multiplexing/Demultiplexing



# Multiplexing/Demultiplexing



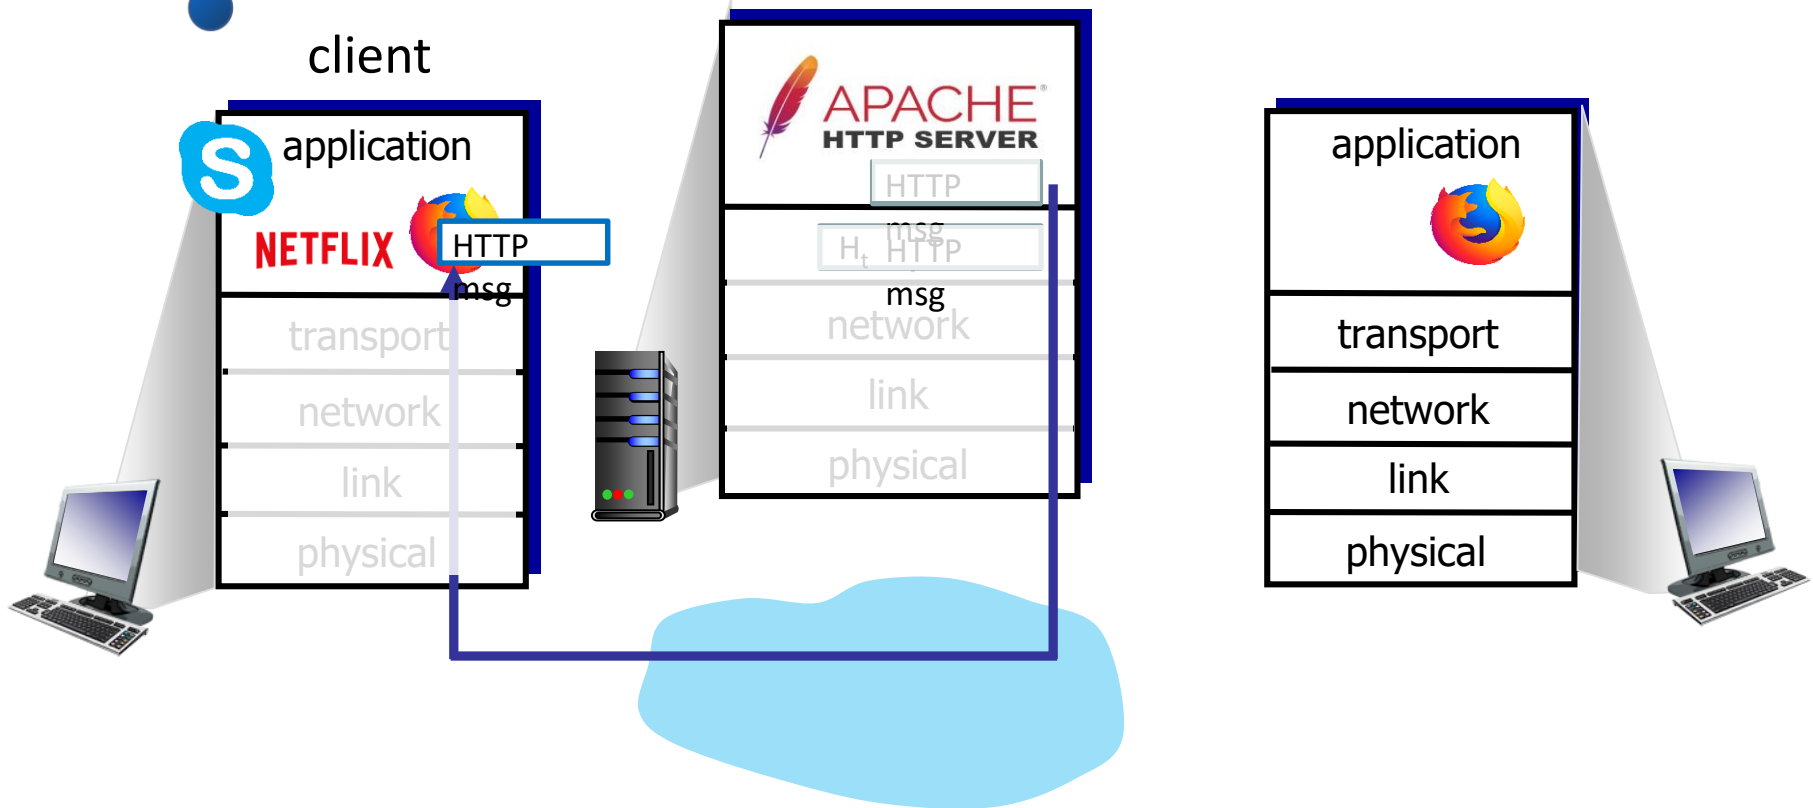
# Multiplexing/Demultiplexing



# Multiplexing/Demultiplexing



*Q: how did transport layer know to deliver message to Firefox browser process rather than Netflix process or Skype process?*



# Multiplexing/Demultiplexing

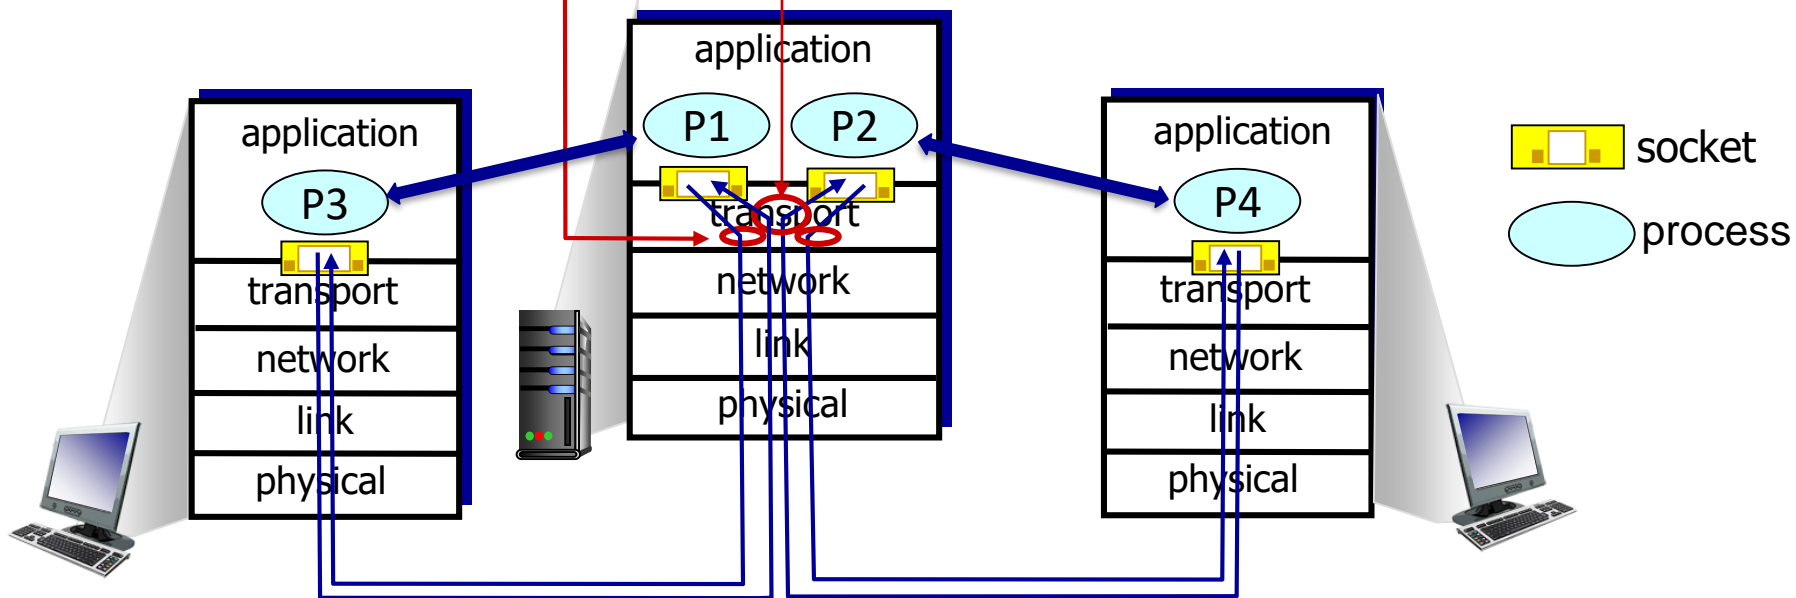


## *multiplexing as sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

## *demultiplexing as receiver:*

use header info to deliver received segments to correct socket



# Multiplexing/Demultiplexing



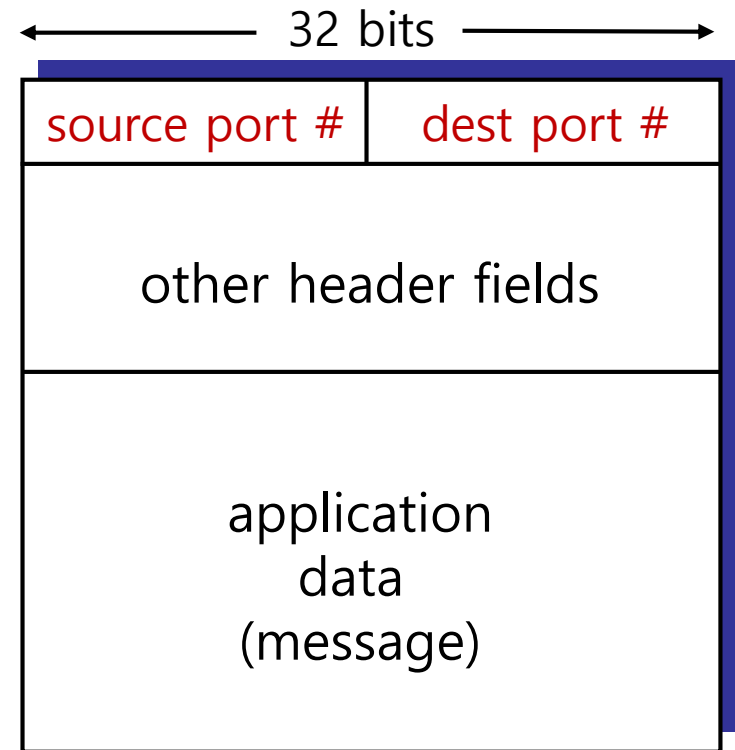
## ■ Multiplexing at send host

- ▶ Gathering data from multiple sockets, enveloping data with header

## ■ Demultiplexing at receive host

- ▶ Delivering received segments to correct socket
- ▶ **Host receives IP datagrams**
  - Each datagram has source IP address, destination IP address
  - Each datagram carries 1 transport-layer segment
  - Each segment has source, destination port number (recall: well-known port numbers for specific applications)

## ■ **Host uses IP addresses & port numbers to direct segment to appropriate socket**



TCP/UDP segment format

# Connectionless Demultiplexing

- Create sockets with port numbers
- UDP socket identified by two-tuple:  
(dest IP address, dest port number)
- When host receives UDP segment:
  - ▶ Checks destination port number in segment
  - ▶ Directs UDP segment to socket with that port number



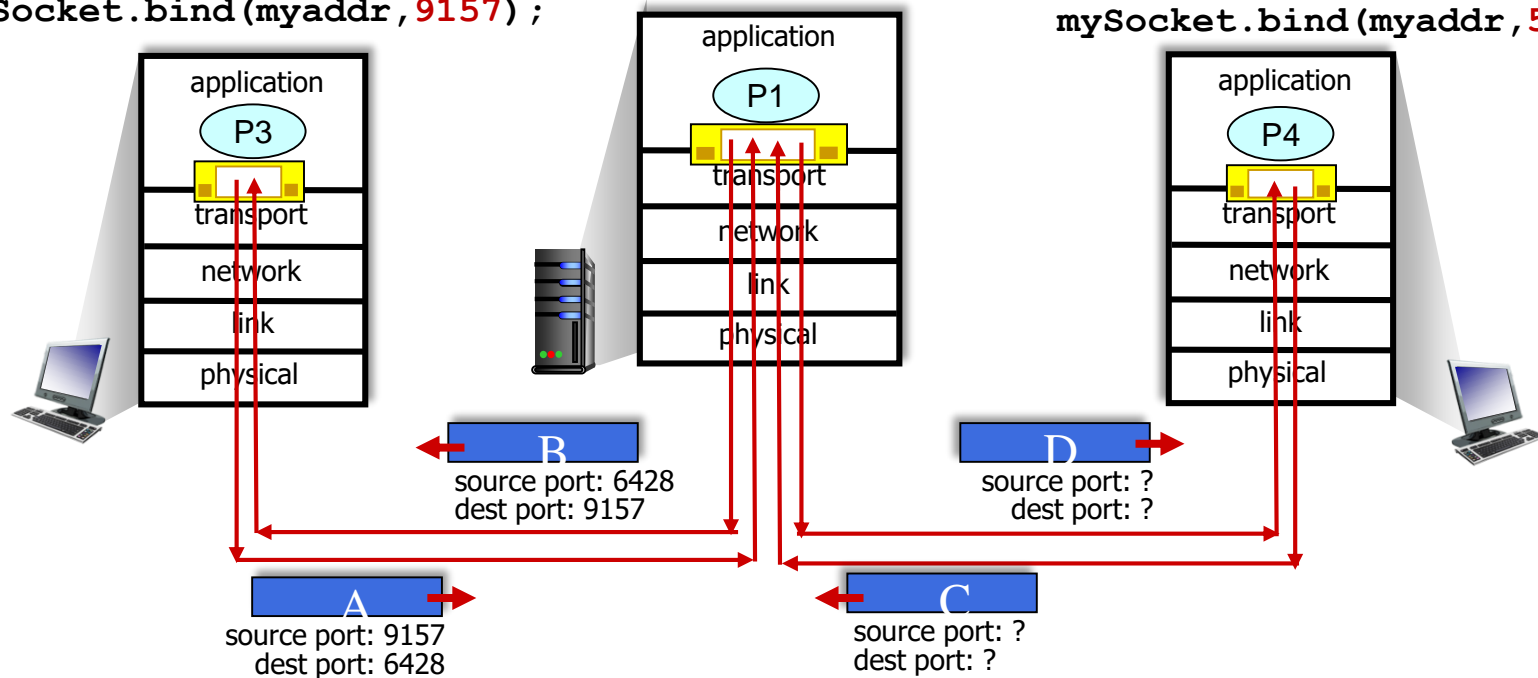
# Connectionless Demultiplexing: an Example



```
mySocket =  
    socket(AF_INET, SOCK_DGRAM)  
mySocket.bind(myaddr, 6428);
```

```
mySocket =  
    socket(AF_INET, SOCK_STREAM)  
mySocket.bind(myaddr, 9157);
```

```
mySocket =  
    socket(AF_INET, SOCK_STREAM)  
mySocket.bind(myaddr, 5775);
```

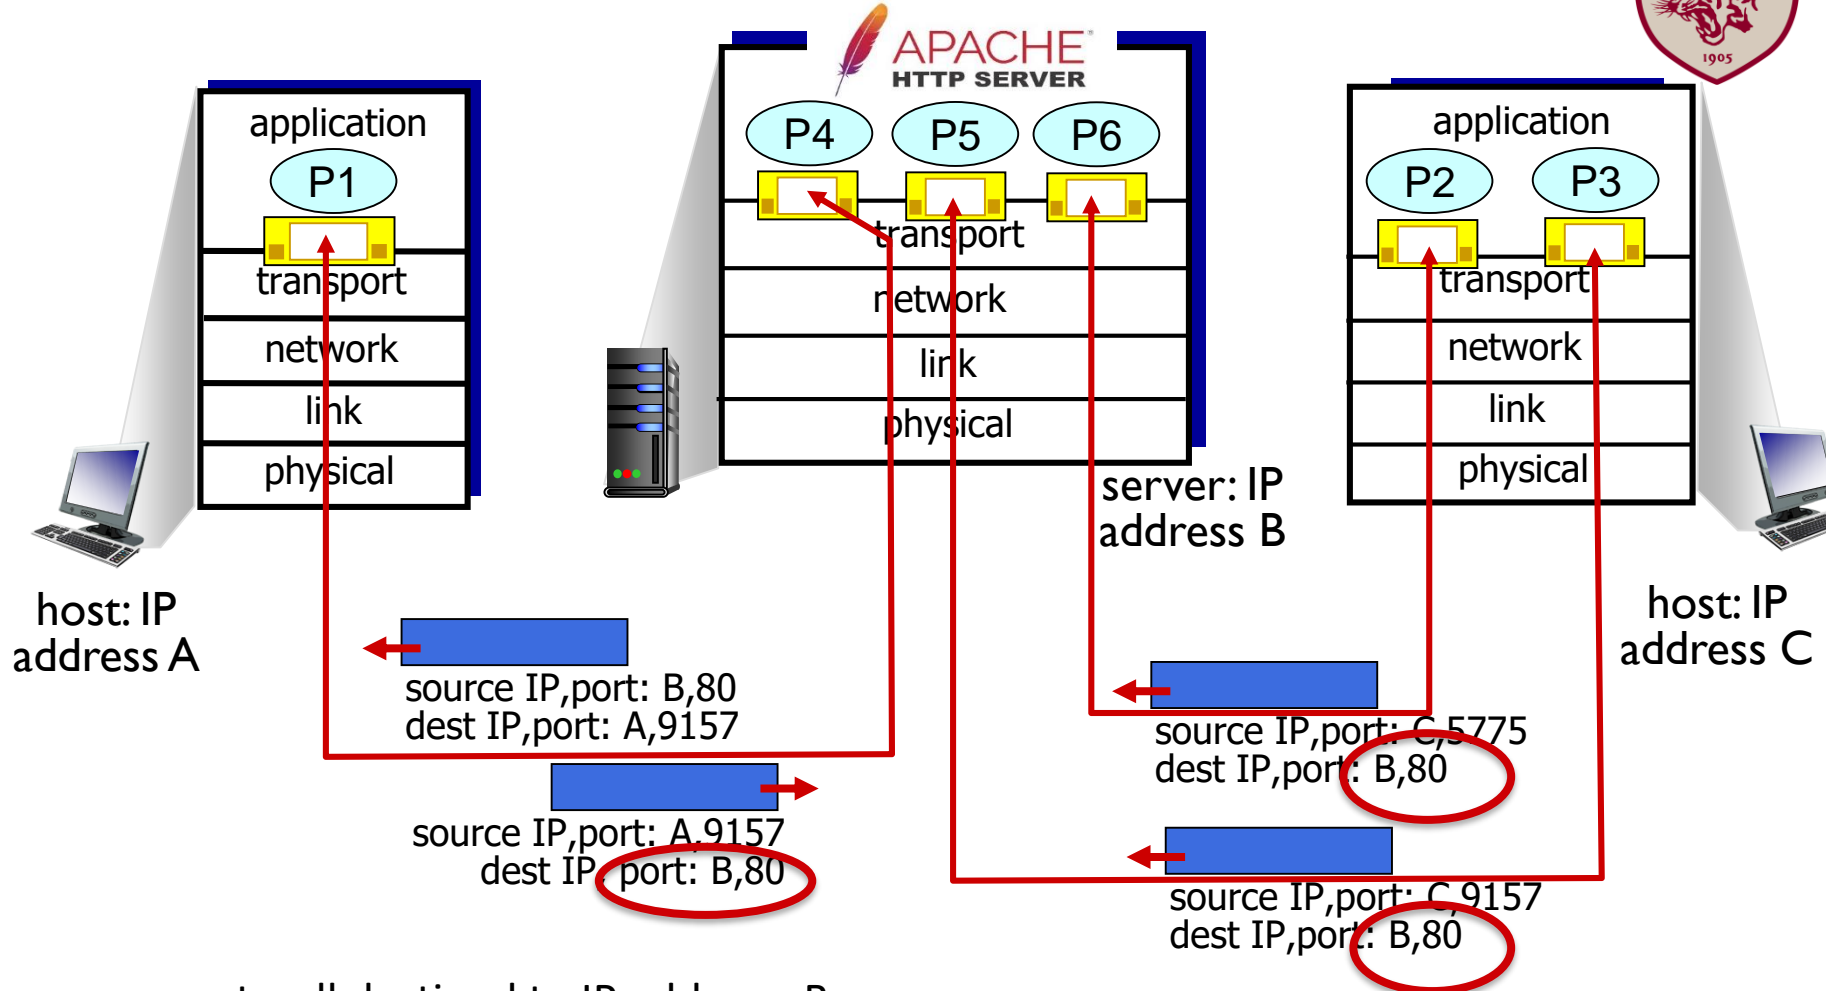


# Connection-Oriented Demultiplexing



- **TCP socket identified by 4-tuple:**
  - ▶ **Source IP address**
  - ▶ **Source port number**
  - ▶ **Destination IP address**
  - ▶ **Destination port number**
- **Receiving host uses all four values to direct segment to appropriate socket**
- **Server host may support many simultaneous TCP sockets:**
  - ▶ Each socket identified by its own 4-tuple
- **Web servers have different sockets for each connecting client**
  - ▶ Non-persistent HTTP will have different socket for each request

# Connection-Oriented Demultiplexing (Cont)



Three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# UDP

# UDP: User Datagram Protocol

- “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - ▶ Lost
  - ▶ Delivered out of order to app
- **Connectionless:**
  - ▶ No handshaking between UDP sender and receiver
  - ▶ Each UDP segment handled independently of others

## Why is there a UDP?

- No connection establishment (which can add delay)
- Simple: no connection state at sender, receiver
- Small segment header
- No congestion control:

# UDP: User Datagram Protocol



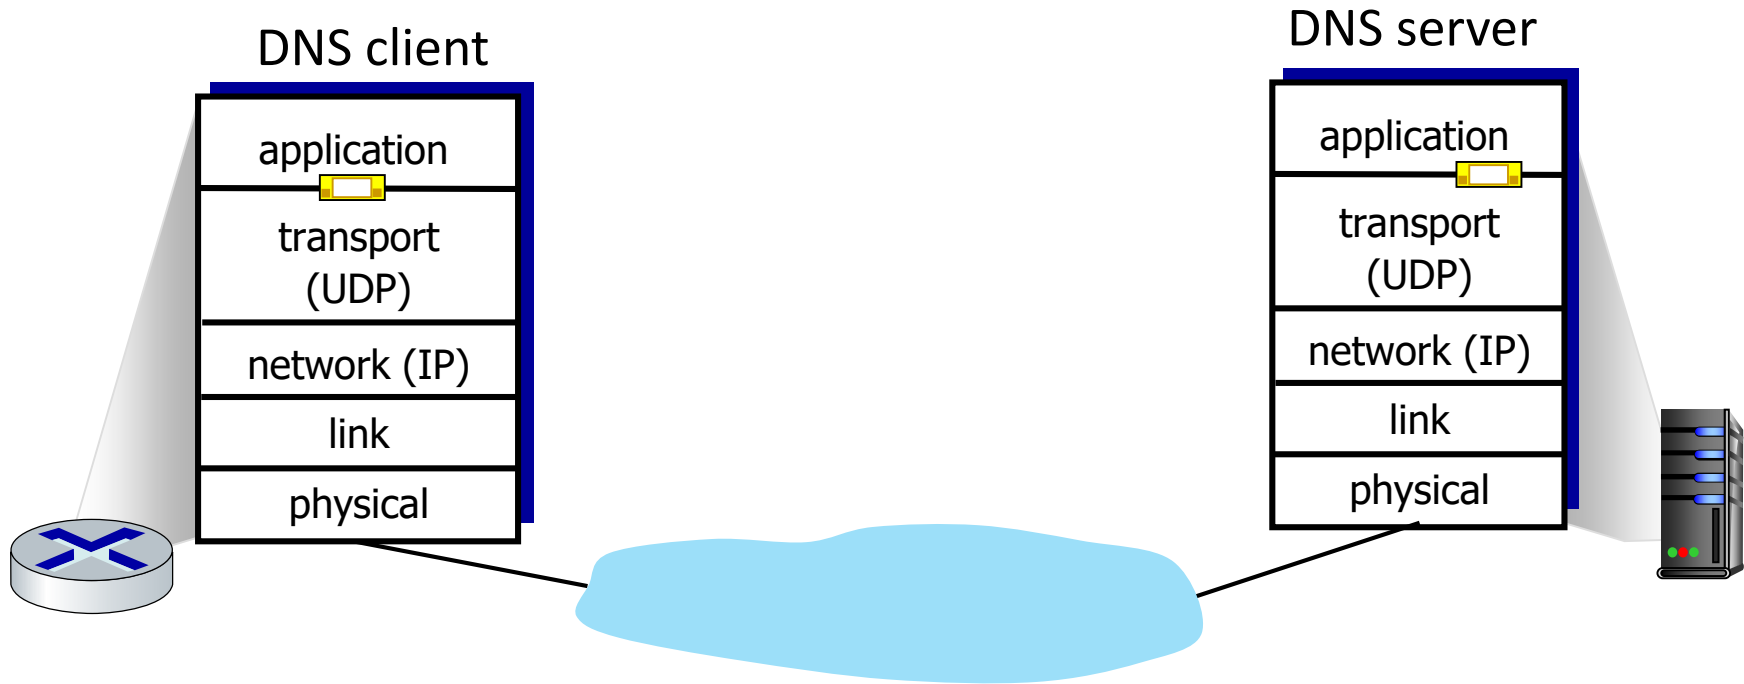
## ■ Other UDP uses

- ▶ DNS
- ▶ SNMP
- ▶ HTTP/3

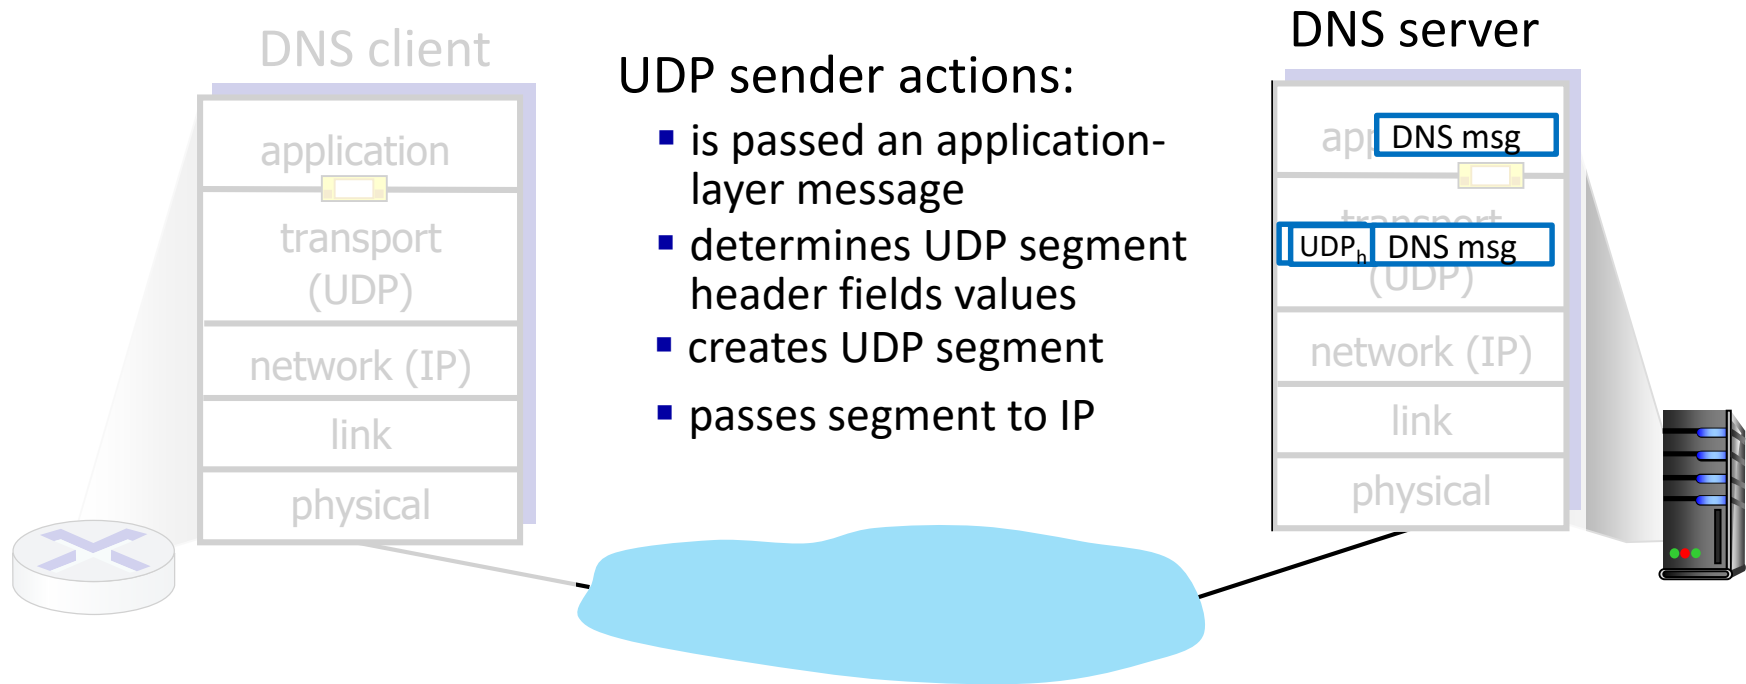
## ■ Reliable transfer over UDP: add reliability at application layer

- ▶ Application-specific error recovery!

# UDP: User Datagram Protocol

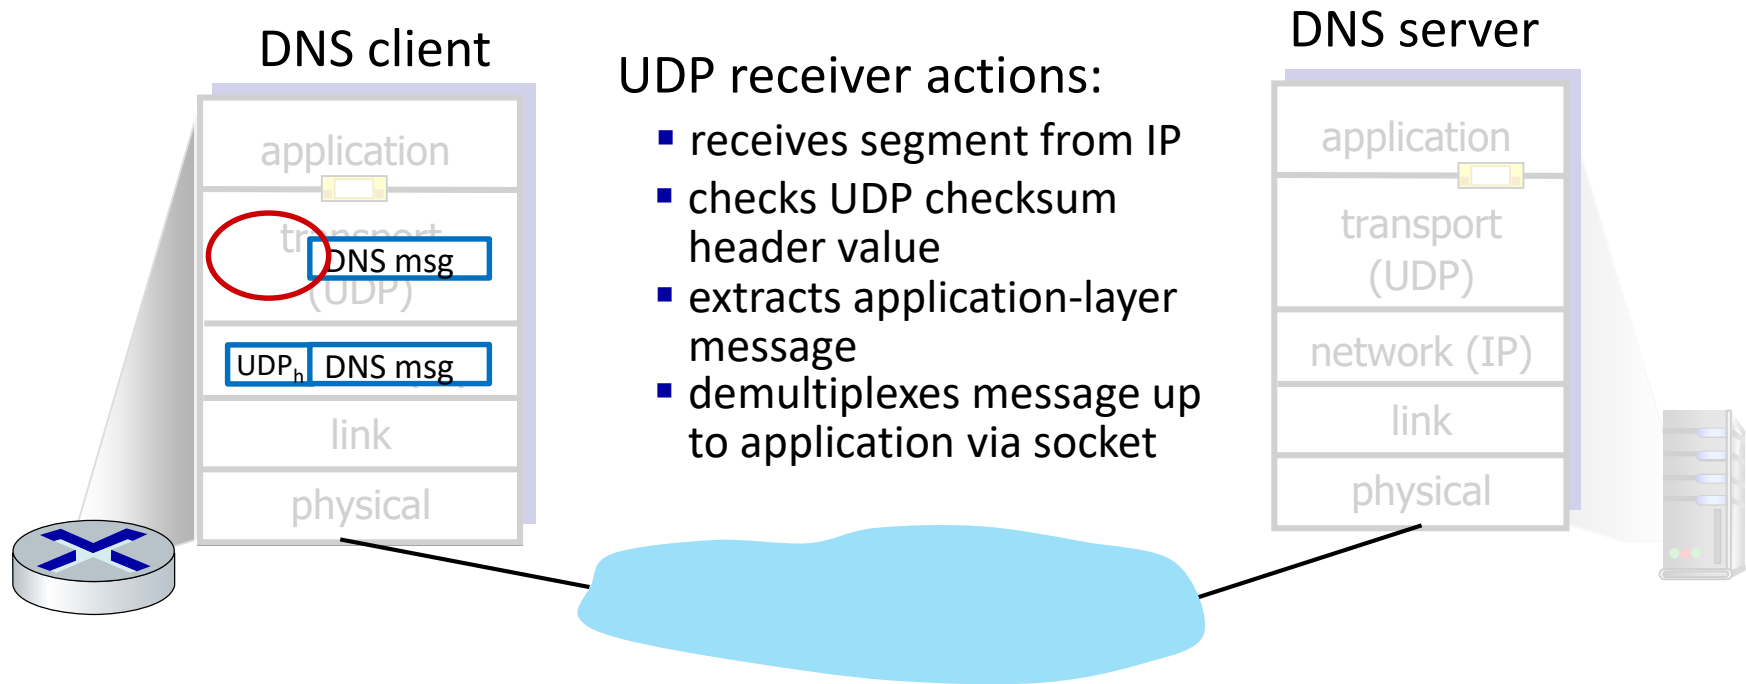


# UDP: User Datagram Protocol

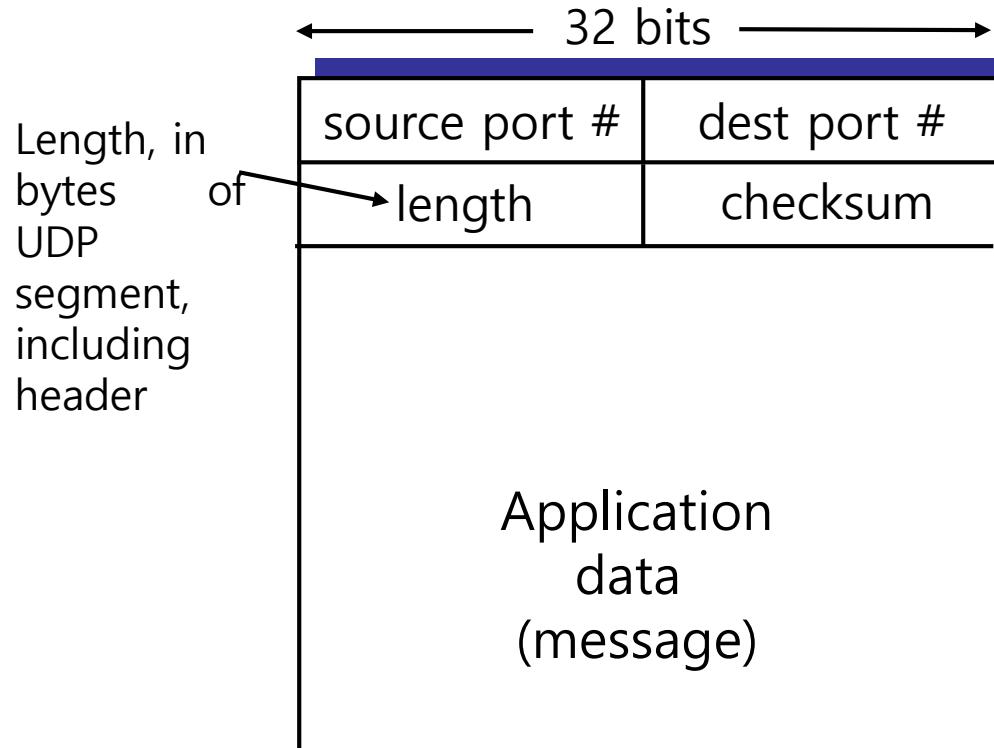




# UDP: User Datagram Protocol



# UDP: User Datagram Protocol



UDP segment format

# UDP Checksum

**Goal:** detect “errors” (e.g., flipped bits) in transmitted segment

## **Sender:**

- Treat segment contents as sequence of 16-bit integers
- Checksum: addition (**1's complement of the sum**) of segment contents
- Sender puts checksum value into UDP checksum field

## **Receiver:**

- Compute checksum of received segment
- Check if computed checksum equals checksum field value:
  - ▶ NO - error detected
  - ▶ YES - no error detected.

# Checksum Example



## ■ Note

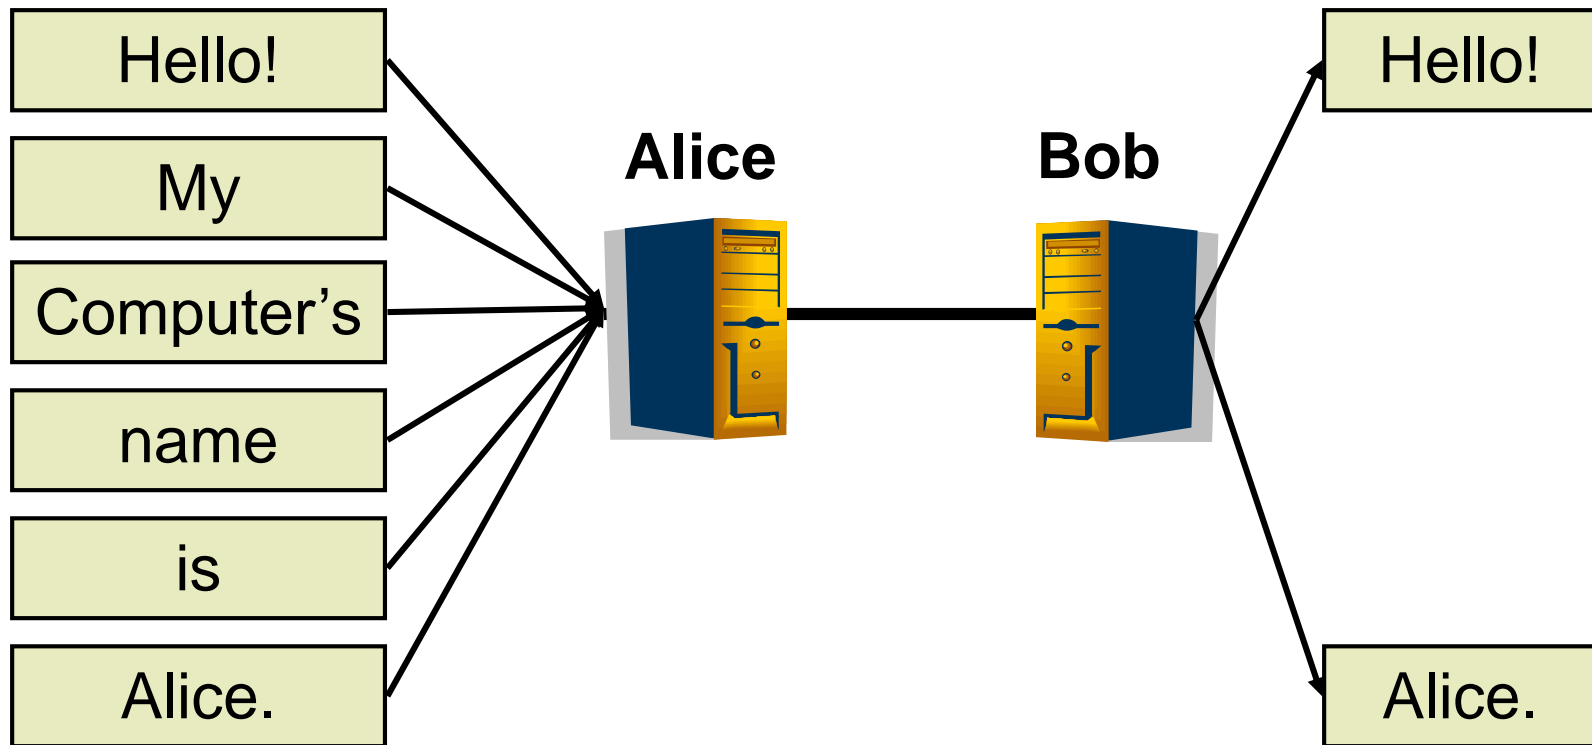
- ▶ When adding numbers, a carryout from the most significant bit needs to be added to the result

## ■ Example: add two 16-bit integers

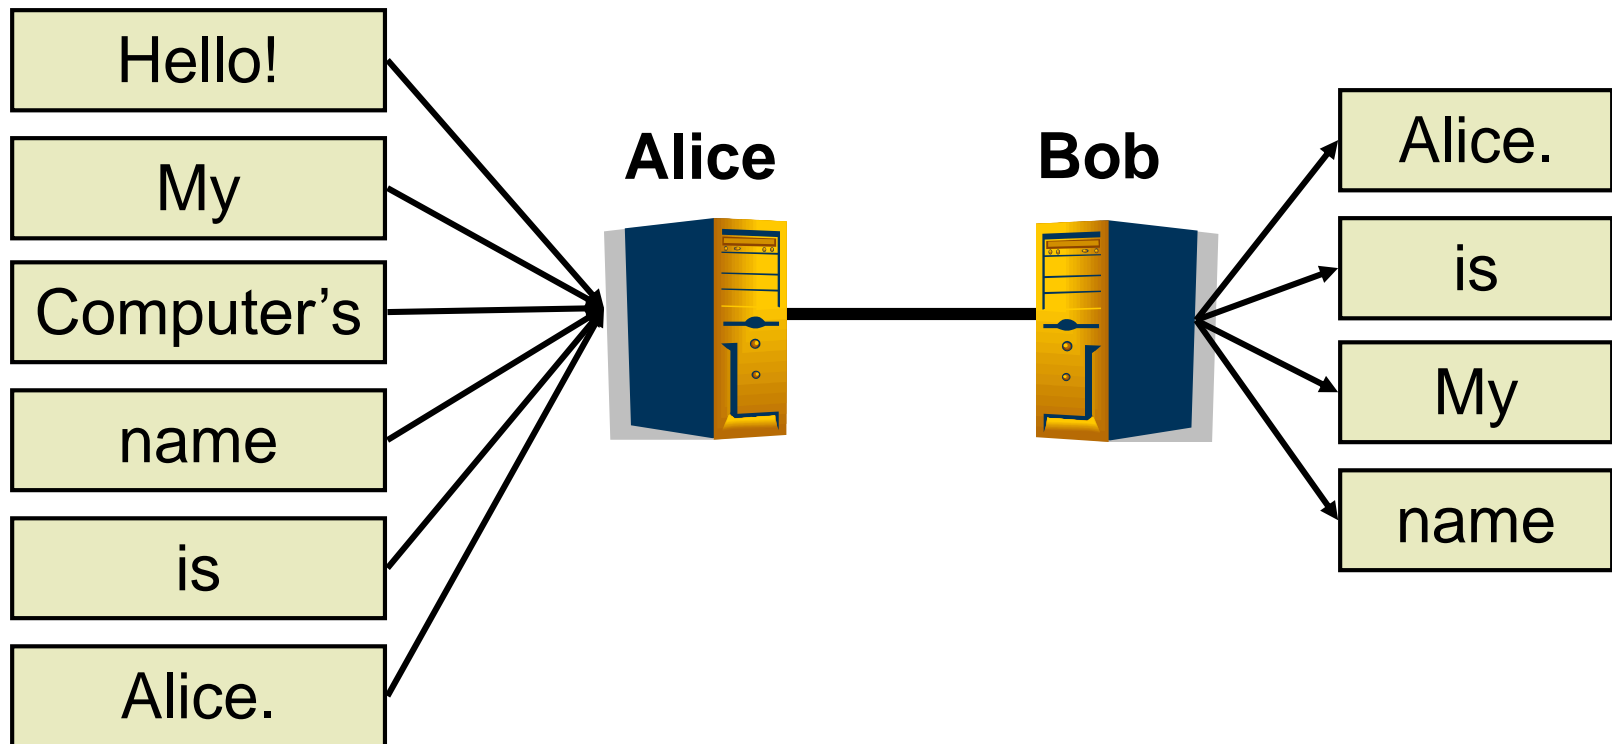
	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

# TCP

# Reliable Transmission



# Reliable Transmission



# Reliable Transmission



- Suppose error protection identifies valid and invalid packets
- Can we make the channel appear reliable?
  - ▶ Insure **packet delivery**
  - ▶ Maintain **packet order**
  - ▶ Provide reliability at full link capacity

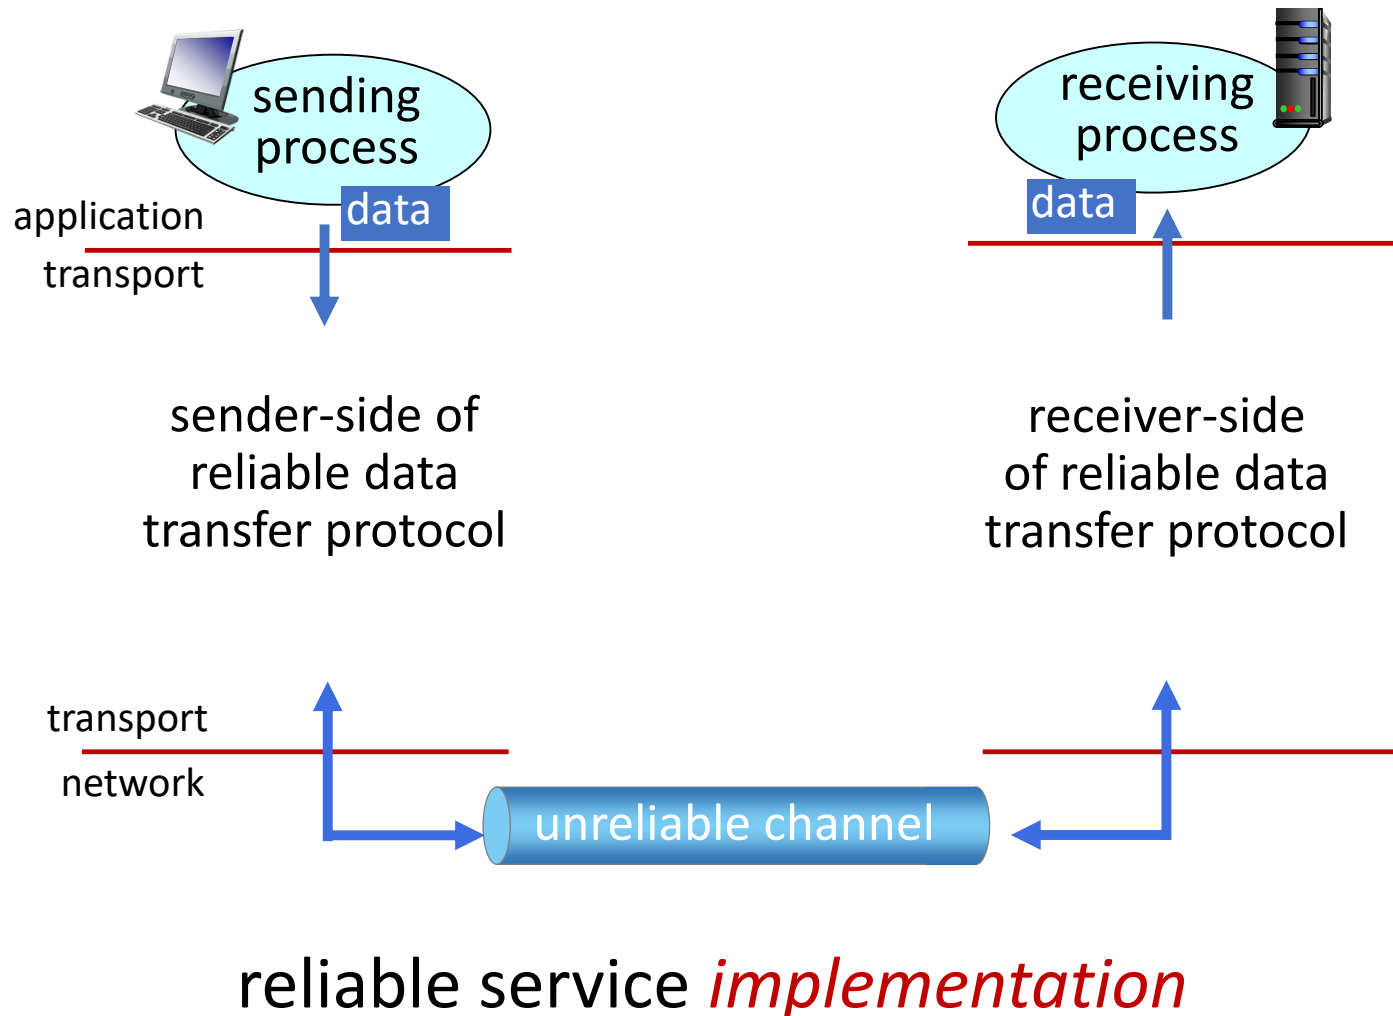


# Reliable Transmission



reliable service *abstraction*

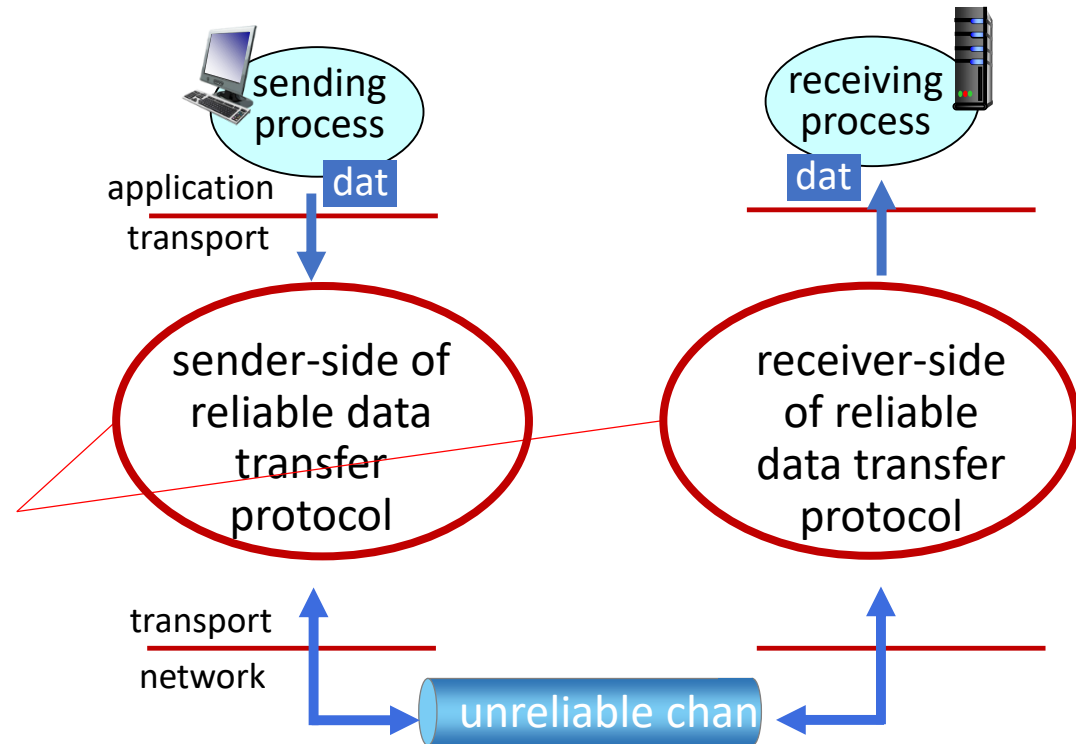
# Reliable Transmission



# Reliable Transmission



Complexity of reliable data transfer protocol will depend (strongly) on characteristics of unreliable channel (lose, corrupt, reorder data?)



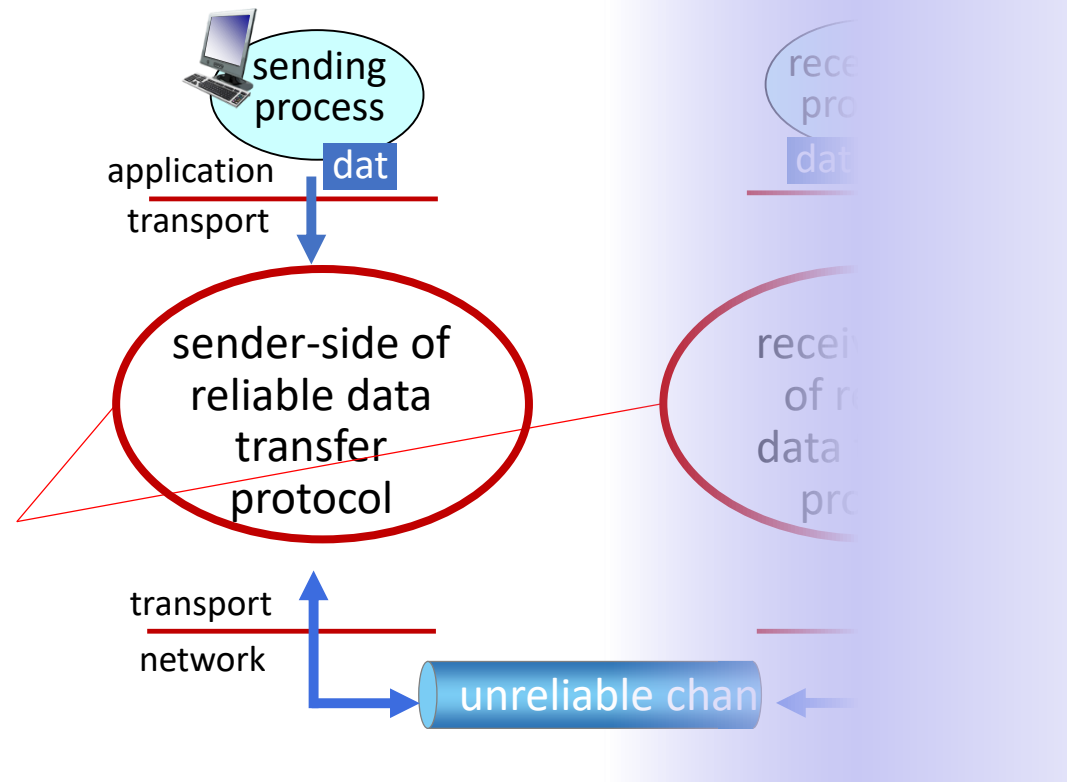
reliable service *implementation*

# Reliable Transmission



Sender, receiver do *not* know the “state” of each other, e.g., was a message received?

- unless communicated via a message



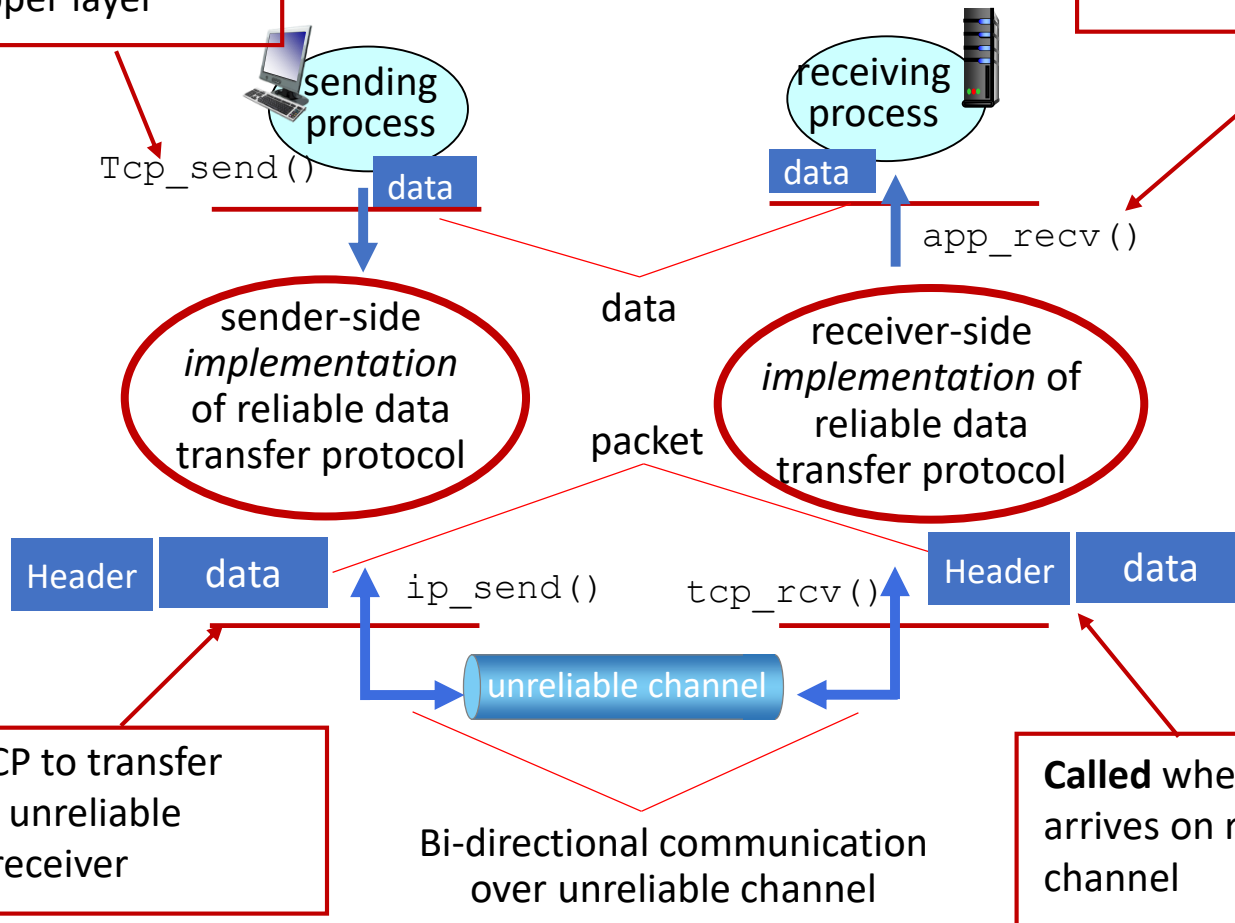
reliable service *implementation*

# Reliable Transmission



**Called** from above, (e.g., by app.). Passed data to deliver to receiver upper layer

**Called** by TCP to deliver data to upper layer



# Reliable Transmission Outline



- **Fundamentals of Automatic Repeat reQuest (ARQ) algorithms**
  - ▶ A family of algorithms that provide reliability through retransmission
- **ARQ algorithms (simple to complex)**
  - ▶ Stop-and-wait
  - ▶ Concurrent logical channels
  - ▶ Sliding window
    - Go-back-n
    - Selective repeat
- **Alternative: forward error correction (FEC)**

## ■ Acknowledgement (**ACK**)

- ▶ Receiver tells the sender when a frame is received
  - Selective acknowledgement (**SACK**)
    - Specifies set of frames received
  - Cumulative acknowledgement (**ACK**)
    - Have received specified frame and all previous
  - Negative acknowledgement (**NAK**)
    - Receiver refuses to accept frame now,  
*e.g.*, when out of buffer space

# Terminology



- **Timeout (TO)**
  - ▶ Sender decides the frame (or ACK) was lost
  - ▶ Sender can try again
- **ARQ also called Positive Acknowledgement with Retransmission (PAR)**



# Stop and Wait

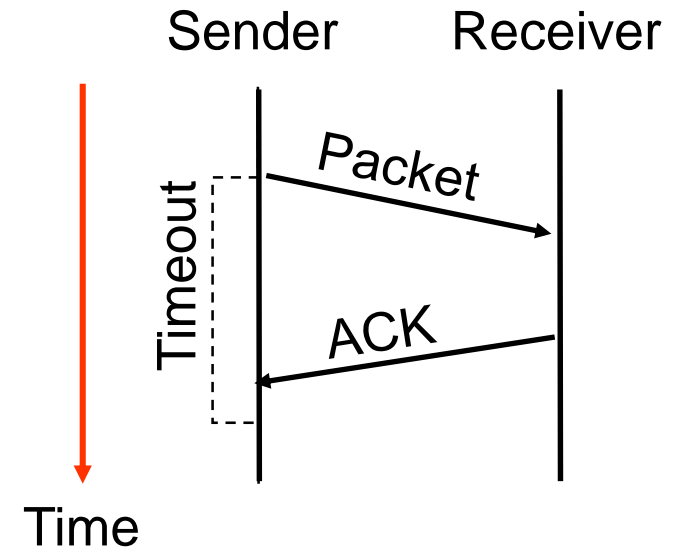


## ■ ARQ

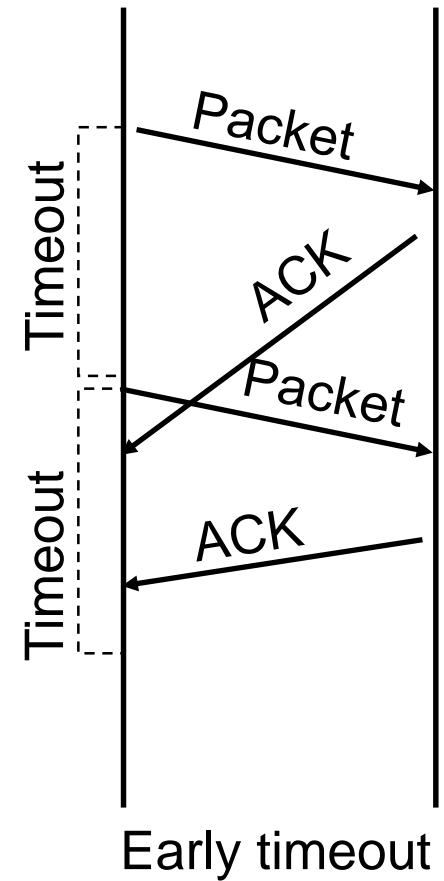
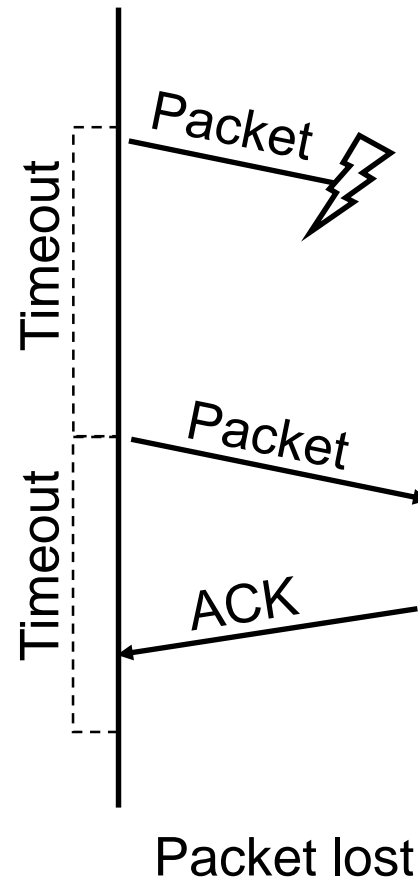
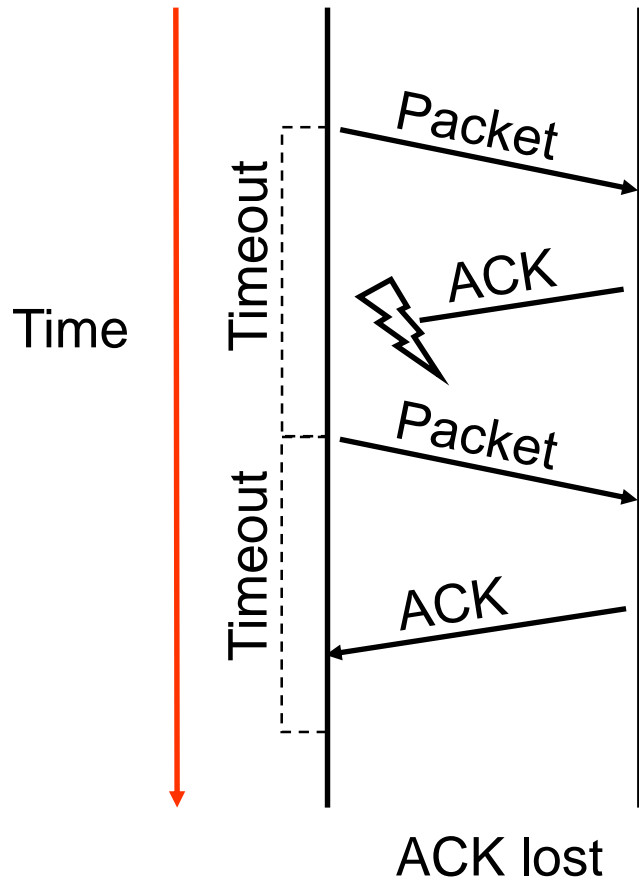
- ▶ Receiver sends acknowledgement (ACK) when it receives packet
- ▶ Sender waits for ACK and timeouts if it does not arrive within some time period

## ■ Simplest ARQ protocol

- Send a packet, stop and wait until ACK arrives

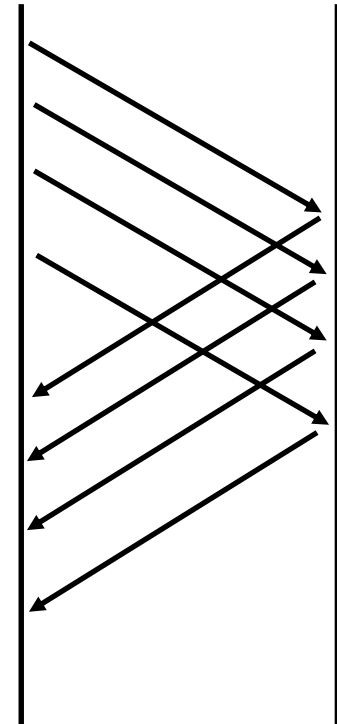


# Stop and Wait



# How to Keep the Pipe Full?

- **Send multiple packets without waiting for first to be acked**
  - ▶ Number of pkts in flight = window
- **Reliable, unordered delivery**
  - ▶ Several parallel stop & waits
  - ▶ Send new packet after each ack
  - ▶ Sender keeps list of unack'ed packets; resends after timeout
  - ▶ Receiver same as stop & wait
- **How large a window is needed?**
  - ▶ Round trip delay \* bandwidth = capacity of pipe

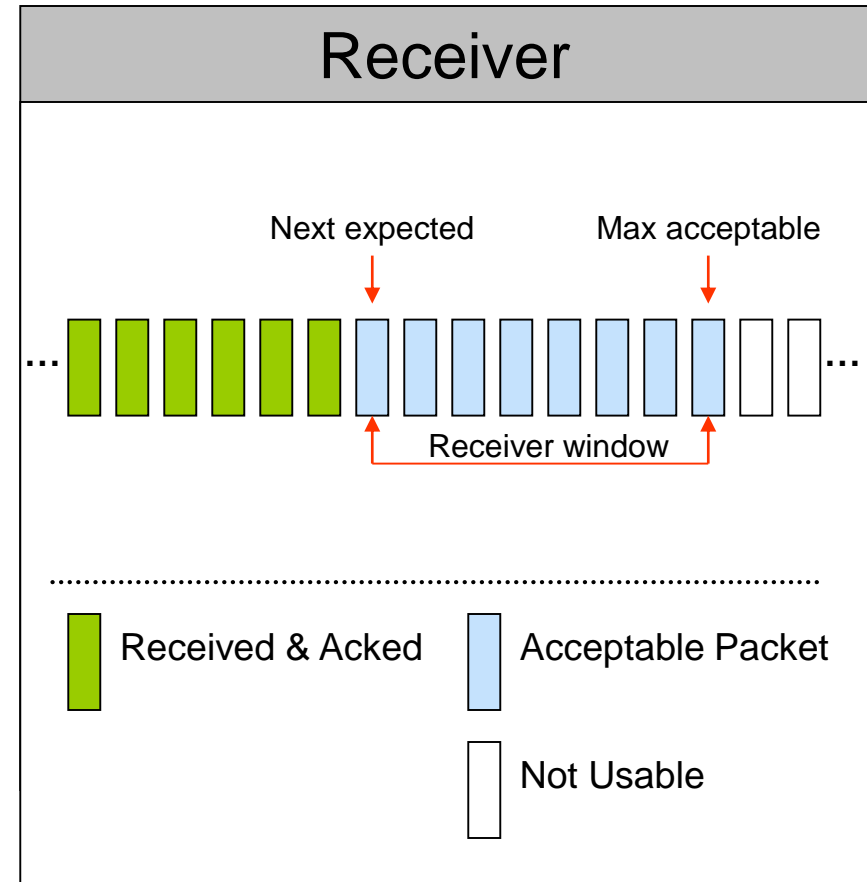
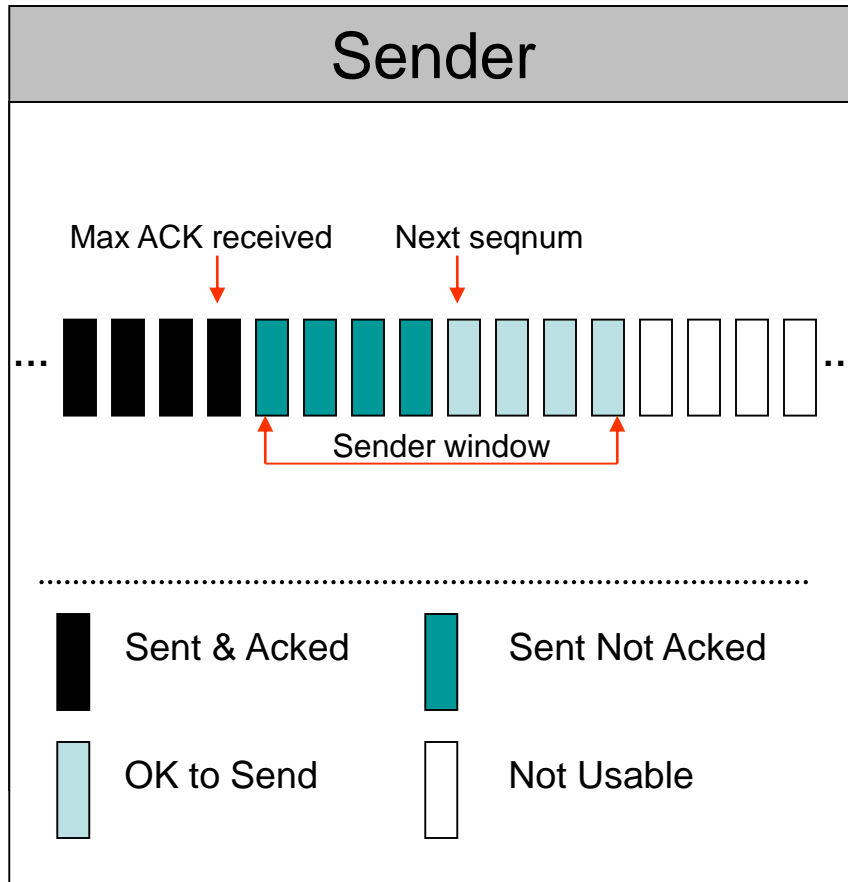


# Sliding Window



- **Reliable, ordered delivery**
- **Receiver has to hold onto a packet until all prior packets have arrived**
  - ▶ Sender must prevent buffer overflow at receiver
- **Circular buffer at sender and receiver**
  - ▶ Packets in transit  $\leq$  buffer size
  - ▶ Advance when sender and receiver agree packets at beginning have been received

# Sliding Window



# Sliding Window: Common Case



- **On reception of new ACK (i.e. ACK for something that was not acked earlier)**
  - ▶ Increase sequence of max ACK received
  - ▶ Send next packet
- **On reception of new in-order data packet (next expected)**
  - ▶ Hand packet to application
  - ▶ Send **cumulative ACK** – acknowledges reception of all packets up to sequence number
  - ▶ Increase sequence of max acceptable packet

# Sliding Window: Loss

- **On reception of out-of-order packet**
  - ▶ Send nothing (wait for source to timeout)
  - ▶ Cumulative ACK (helps source identify loss)
- **Timeout (Go-Back-N recovery)**
  - ▶ Set timer upon transmission of packet
  - ▶ **Retransmit all unacknowledged packets**
- **Performance during loss recovery**
  - ▶ No longer have an entire window in transit
  - ▶ **Can have much more clever loss recovery**

The diagram illustrates a Stop-and-Wait protocol scenario. The **Sender** (top) and **Receiver** (bottom) are shown. A horizontal timeline at the bottom indicates the progression of time.

- Initial State:** The Sender has a buffer containing messages 0 through 10. The Receiver has a buffer containing messages 0 and 1.
- First Round:**
  - The Sender sends message 0 (solid arrow).
  - The Receiver receives message 0 (solid arrow) and immediately sends back **ACK 0** (dashed arrow).
  - The Sender receives **ACK 0** and immediately sends message 1 (solid arrow).
  - The Receiver receives message 1 (solid arrow) and immediately sends back **ACK 1** (dashed arrow).
- Error and Timeout:**
  - After receiving **ACK 1**, the Sender attempts to send message 2 (solid arrow).
  - However, a red arrow labeled **E** (Error) points to the message 2 in the Sender's buffer.
  - A horizontal double-headed arrow labeled **timeout** spans from the point where message 2 was sent to the point where the Sender starts sending message 2 again.
  - During this timeout, the Receiver discards messages 2 through 8, indicated by dashed arrows pointing to 'D' in the Receiver's buffer.
- Retransmission:**
  - After the timeout expires, the Sender resends message 2 (solid arrow).
  - The Receiver receives message 2 (solid arrow) and immediately sends back **ACK 2** (dashed arrow).
  - The Sender receives **ACK 2** and immediately sends message 3 (solid arrow).
  - The Receiver receives message 3 (solid arrow) and immediately sends back **ACK 3** (dashed arrow).
  - This pattern continues for messages 4 through 10, with each message being received and acknowledged immediately.



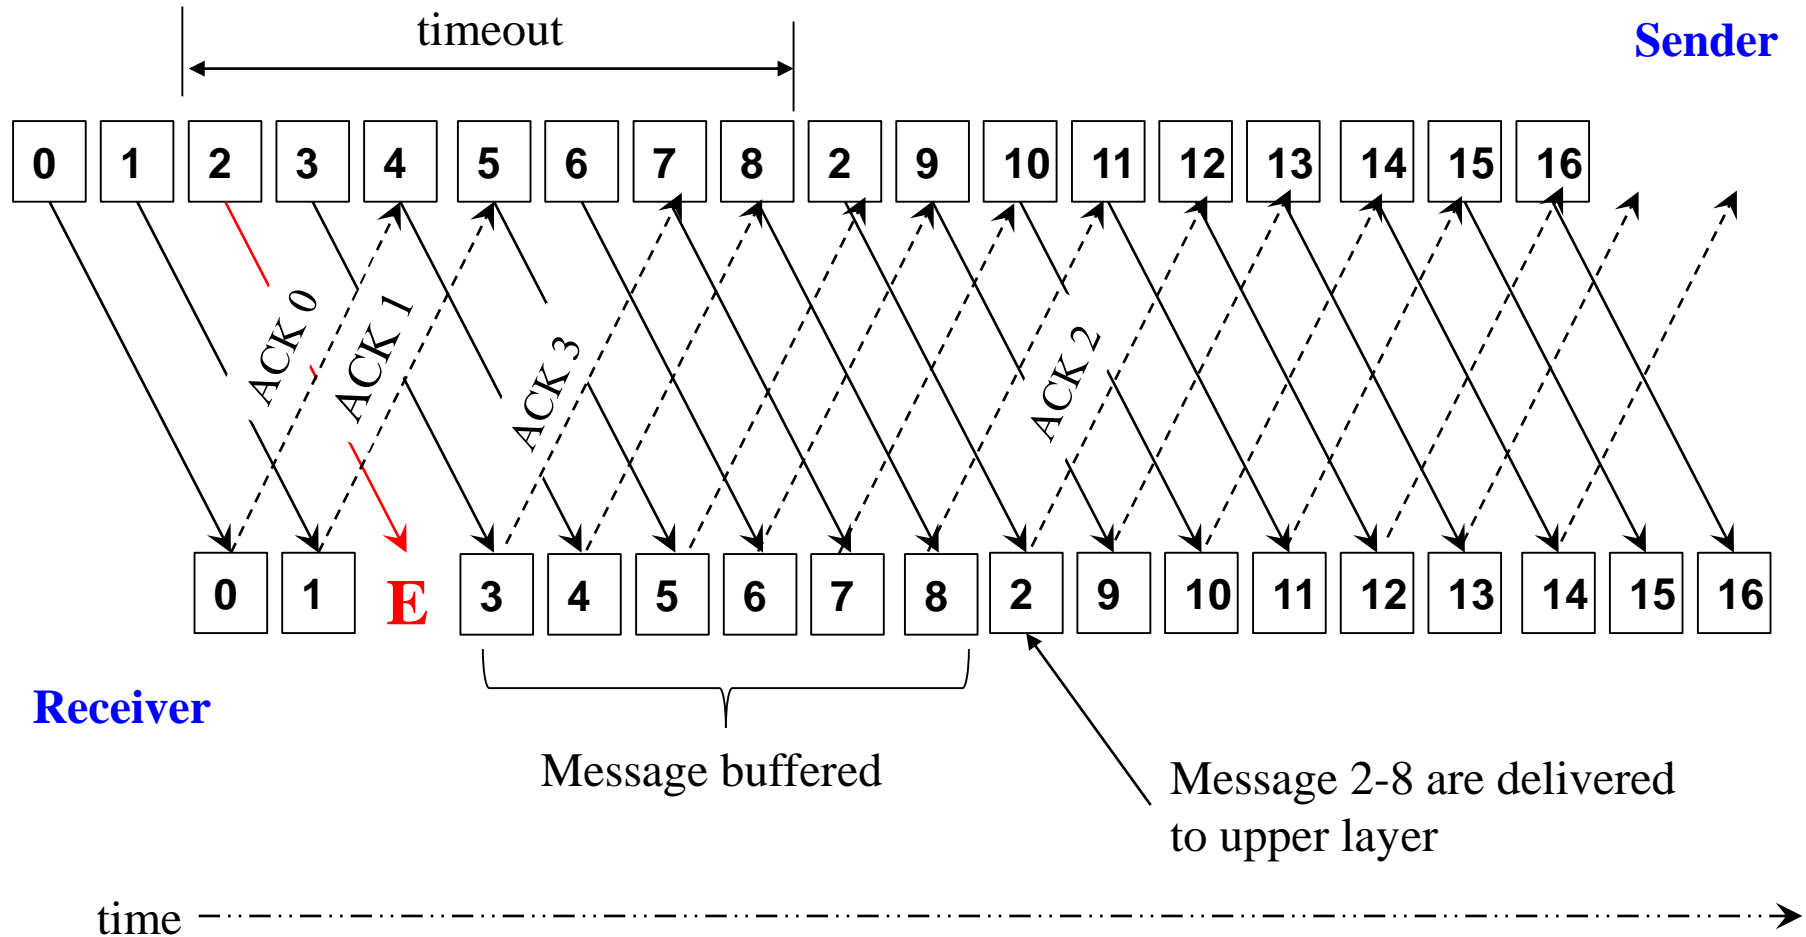
# Sliding Window: Selective Repeat

- **Receiver individually acknowledges all correctly received packets**
  - ▶ Buffers packets, as needed, for **eventual in-order delivery** to upper layer
- **Sender only resends packets for which ACK not received**
  - ▶ Sender timer for each unACKed packet

# Sliding Window: Go-Back-N



Sender



Receiver

Message buffered

Message 2-8 are delivered to upper layer

time

# Sliding Window: Sequence Number

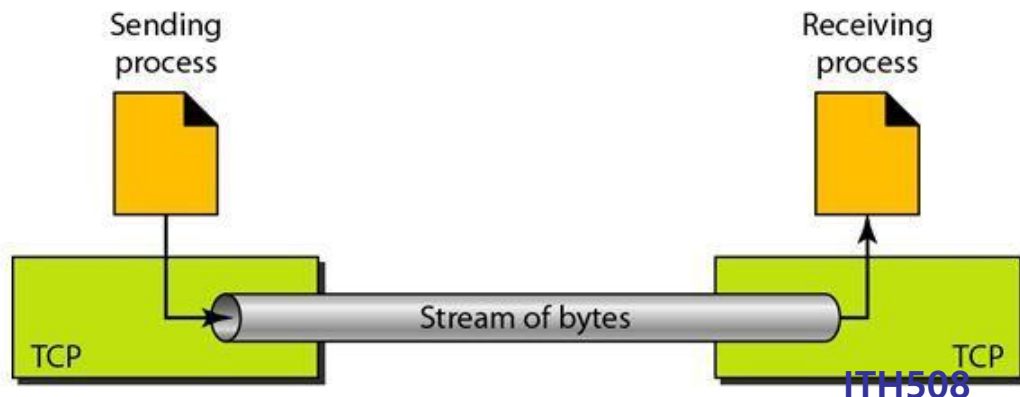


- **How large do sequence numbers need to be?**
  - ▶ Must be able to detect wrap-around
  - ▶ Depends on sender/receiver window size
- **E.g.**
  - ▶ Max seq = 7, send win=recv win=7
  - ▶ If pkts 0..6 are sent successfully and all acks lost
    - Receiver expects 7, 0..5, sender retransmits old 0..6!!!
- **Max sequence must be  $\geq$  send window + recv window**

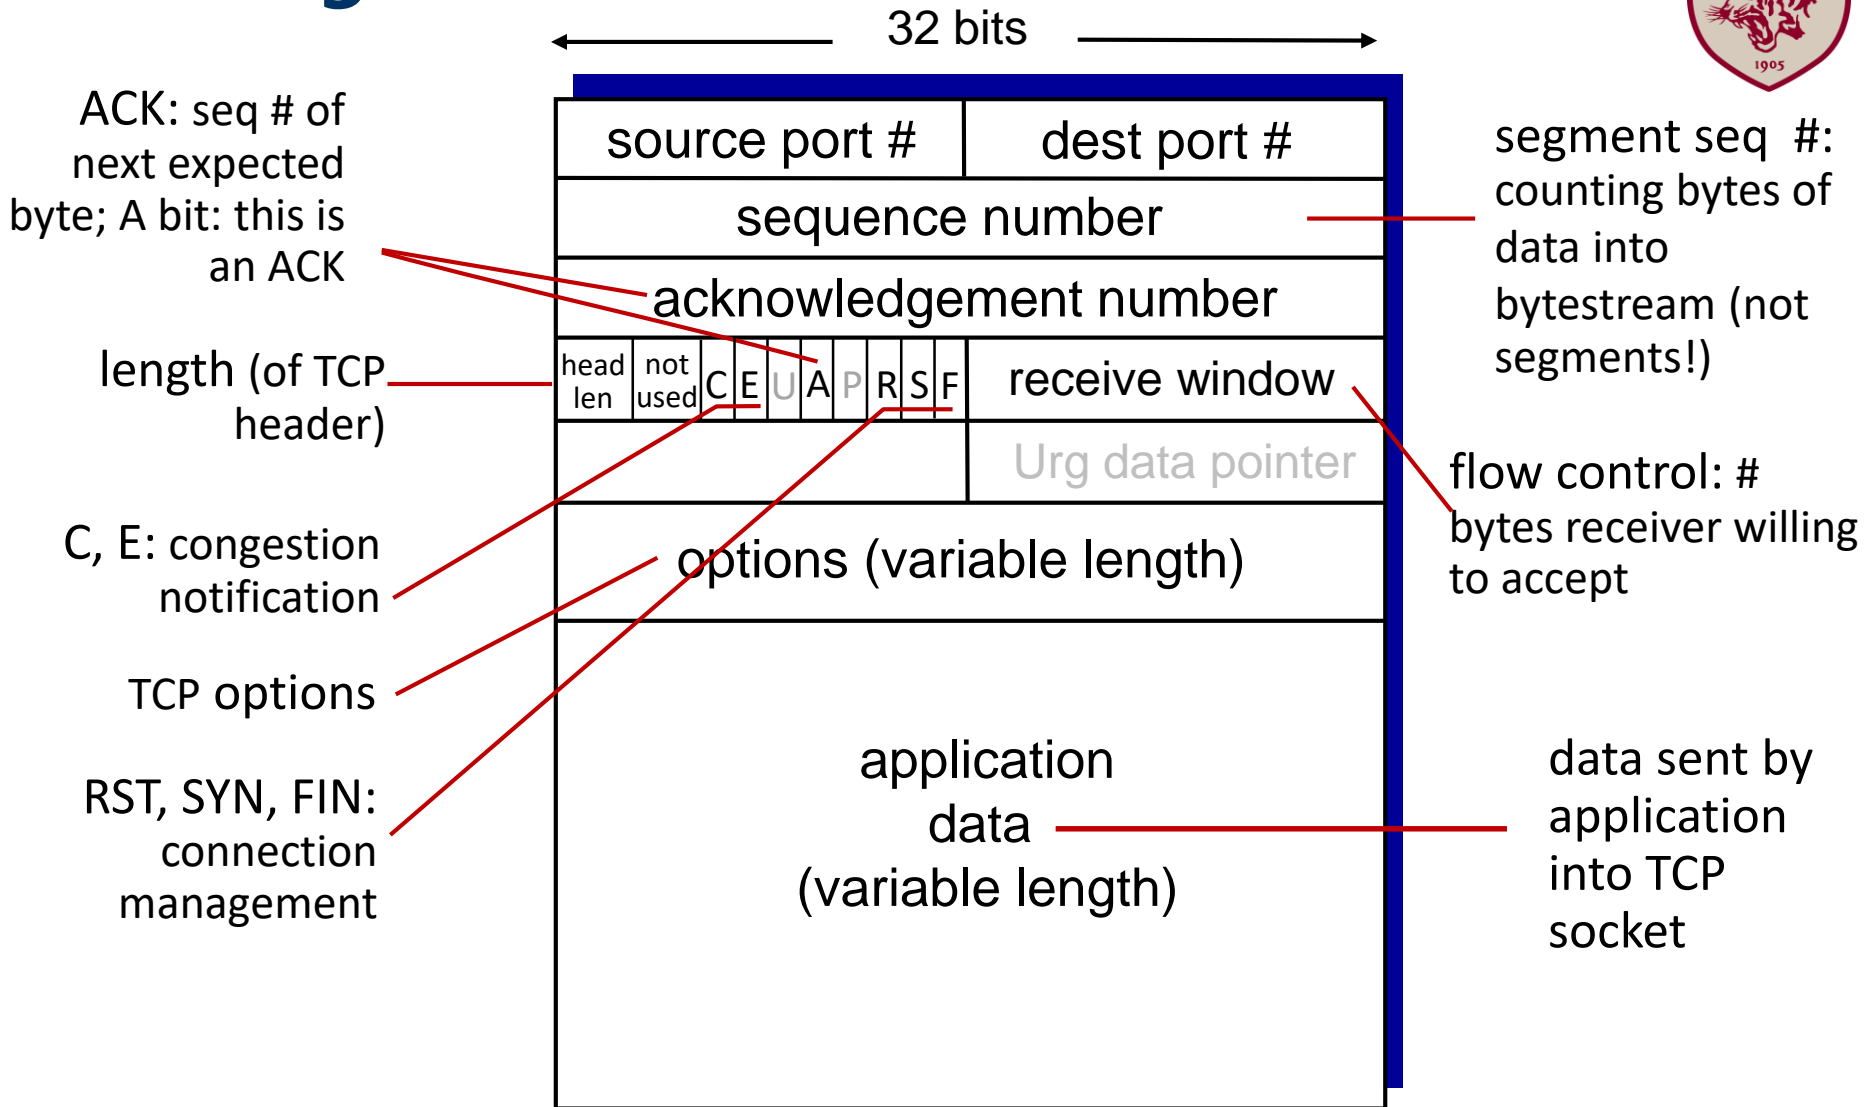
# TCP DYANAMICS

# TCP: Overview

- **Point-to-point:**
  - ▶ One sender, one receiver
- **Reliable, in-order *byte stream*:**
  - ▶ No "message boundaries"
- **Pipelined:**
  - ▶ TCP congestion and flow control set window size
- ***Send & receive buffers***
- **Full duplex data:**
  - ▶ Bi-directional data flow in same connection
  - ▶ MSS: maximum segment size
- **Connection-oriented:**
  - ▶ Handshaking (exchange of control messages) init's sender, receiver state before data exchange
- **Flow controlled:**
  - ▶ Sender will not overwhelm receiver



# TCP Segment Structure



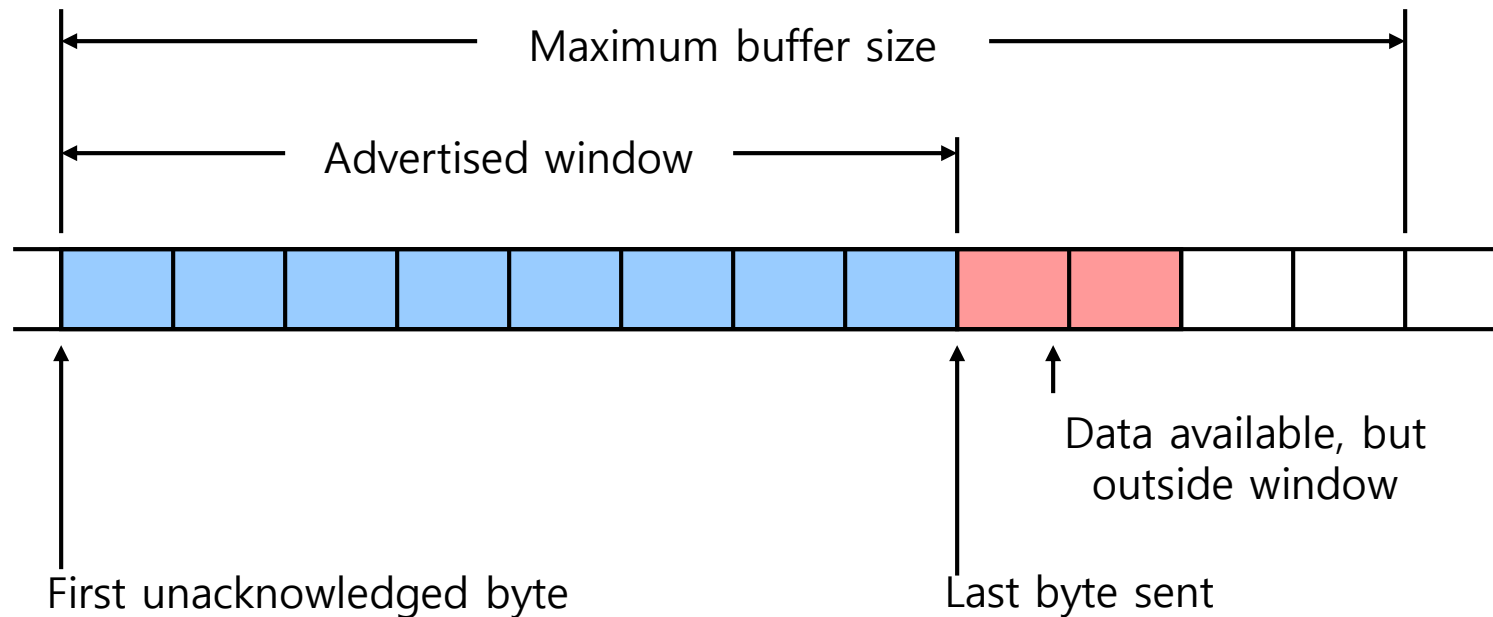
# TCP Sliding Window Protocol

- **Sequence numbers**
  - ▶ Indices into byte stream
- **ACK sequence number**
  - ▶ Actually next byte expected as opposed to last byte received
- **Advertised window**
  - ▶ Enables dynamic receive window size
- **Receive buffers**
  - ▶ Data ready for delivery to application until requested
  - ▶ Out-of-order data out to maximum buffer capacity
- **Sender buffers**
  - ▶ Unacknowledged data
  - ▶ Unsent data out to maximum buffer capacity

# TCP Sliding Window Protocol: Sender Side



- `LastByteAcked <= LastByteSent`
- `LastByteSent <= LastByteWritten`
- Buffer bytes between `LastByteAcked` and `LastByteWritten`

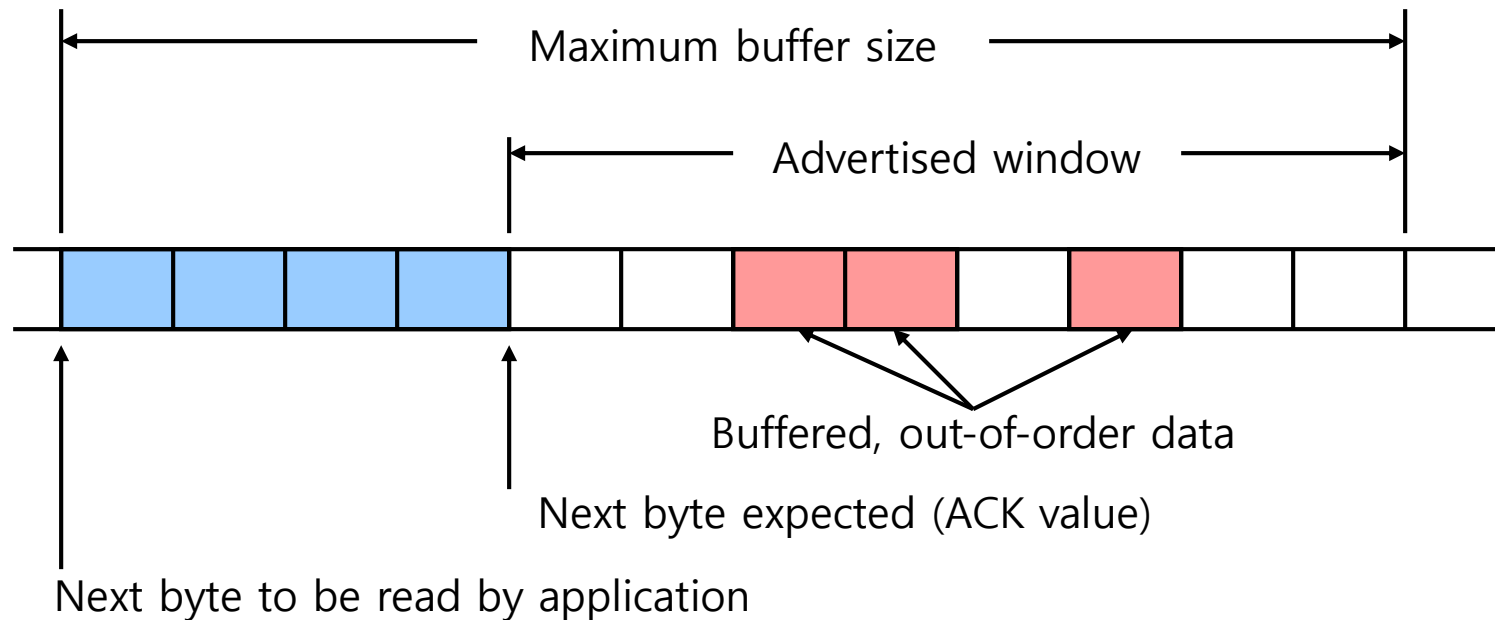




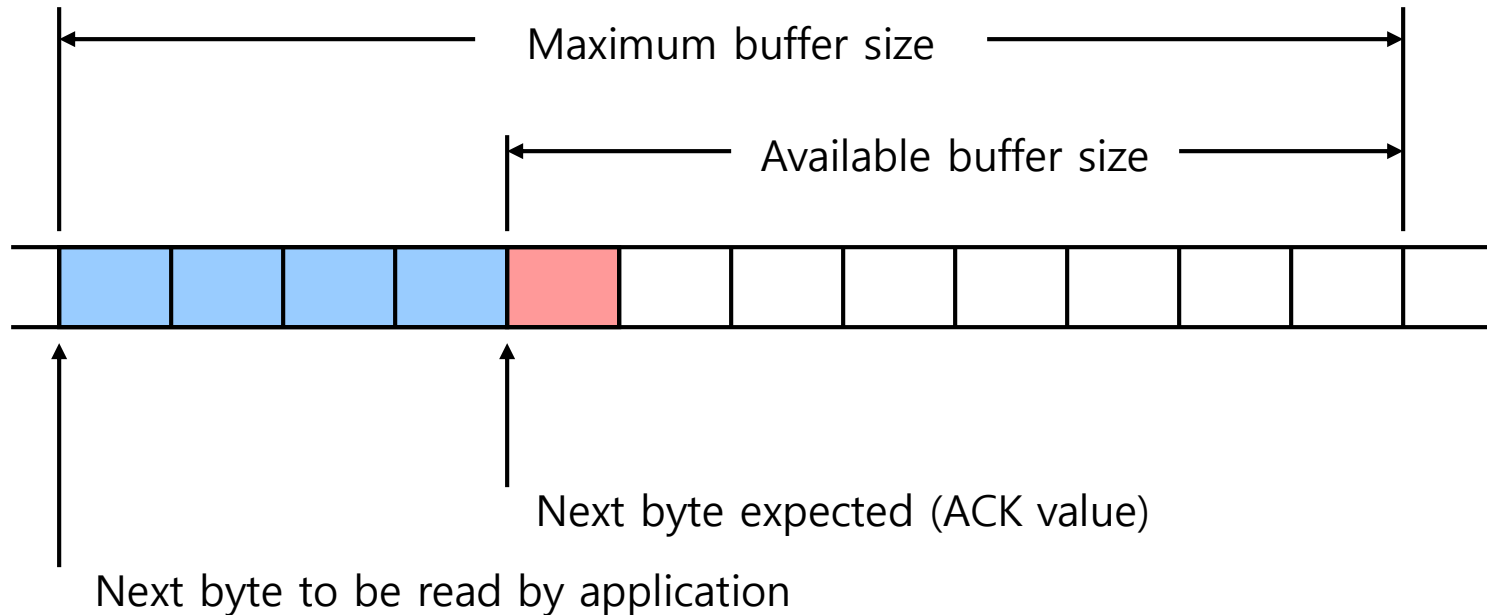
# TCP Sliding Window Protocol: Receiver Side



- $\text{LastByteRead} < \text{NextByteExpected}$
- $\text{NextByteExpected} \leq \text{LastByteRcvd} + 1$
- Buffer bytes between  $\text{NextByteRead}$  and  $\text{LastByteRcvd}$

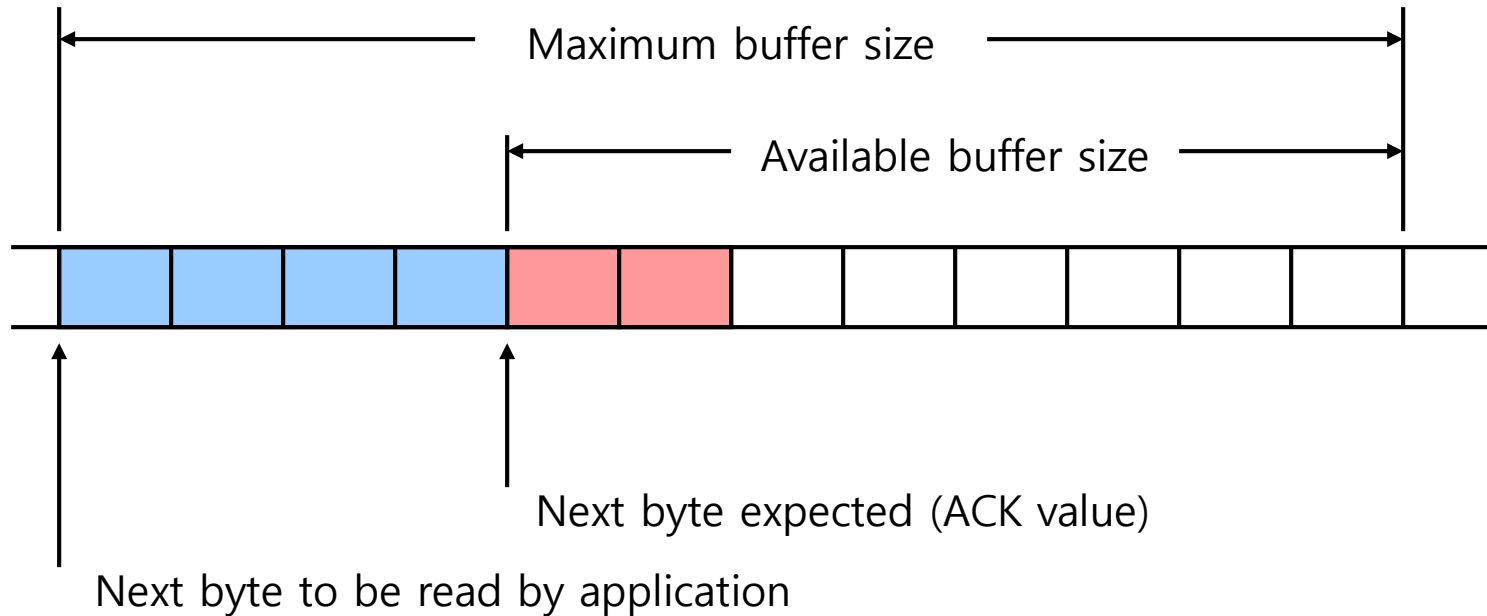


# TCP ACK Generation - 1



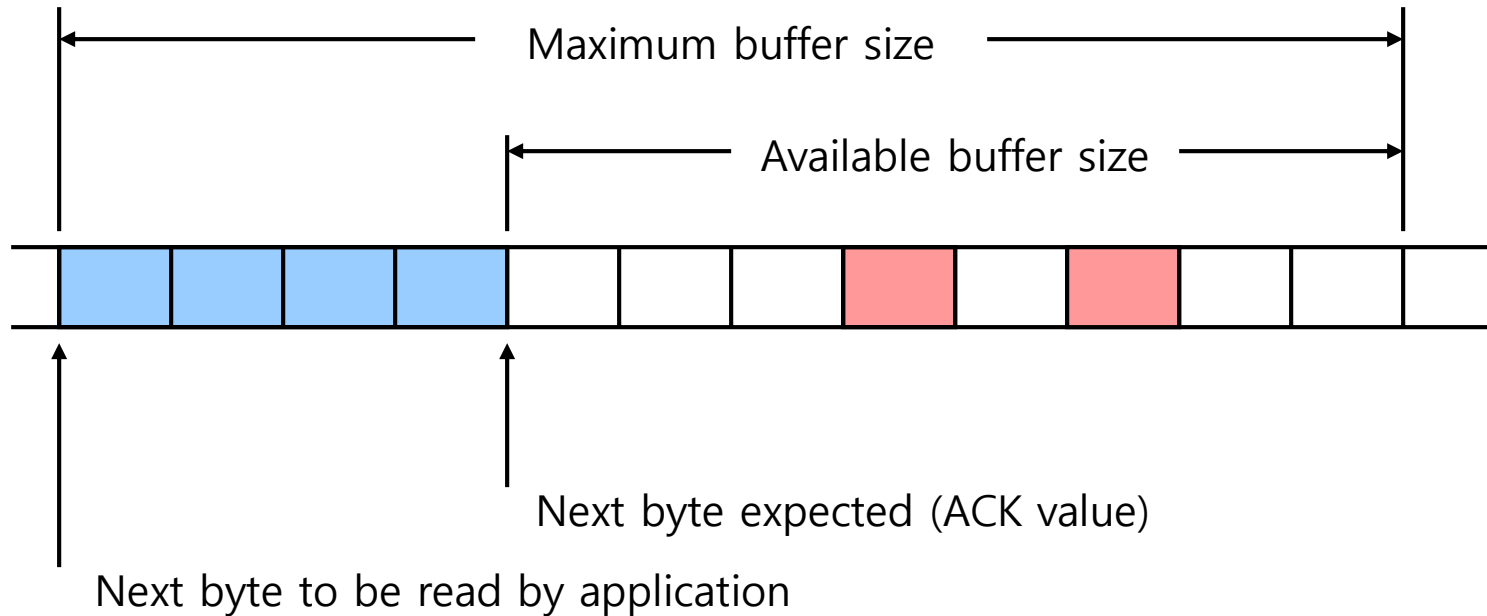
- **Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed**
  - ▶ Delayed ACK. Wait up to 500ms for next segment.

# TCP ACK Generation - 2



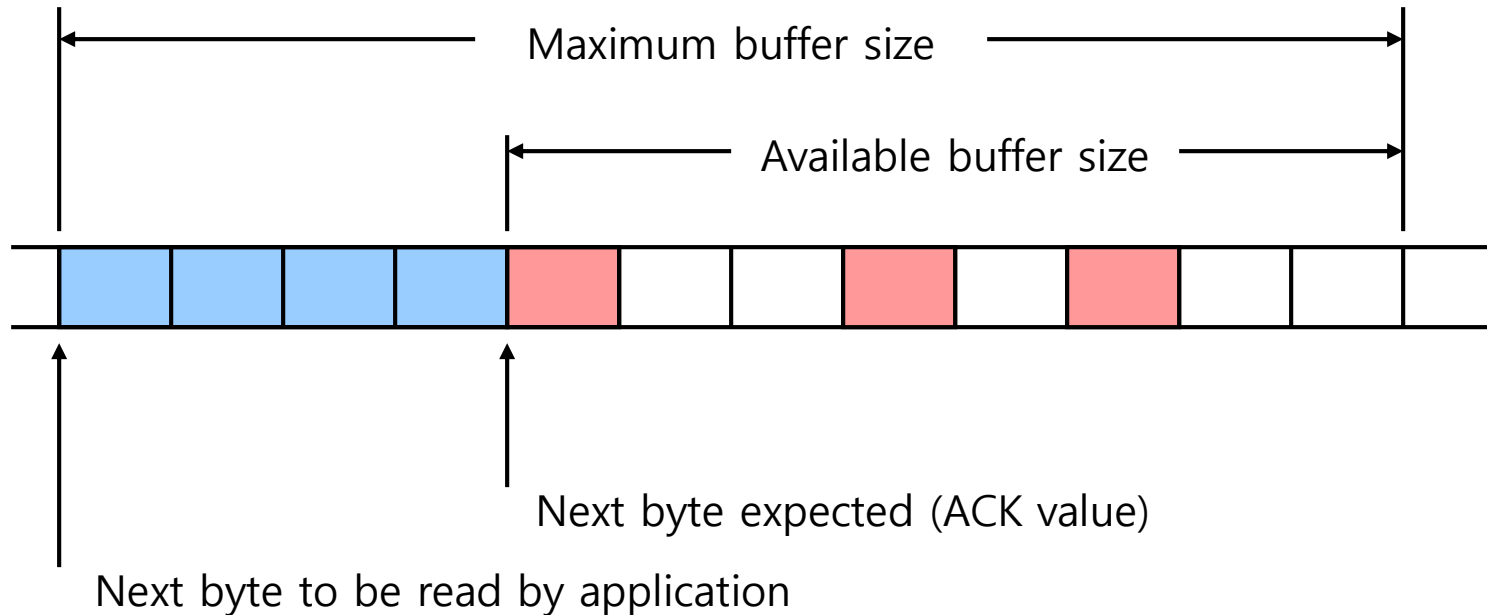
- **Arrival of in-order segment with expected seq #. One other segment has ACK pending**
  - ▶ Immediately send single cumulative ACK, ACKing both in-order segments

# TCP ACK Generation - 3



- **Arrival of out-of-order segment higher-than-expect seq. # Gap detected**
  - ▶ Immediately send duplicate ACK, indicating seq. # of next expected byte

# TCP ACK Generation - 4



- **Arrival of segment that partially or completely fills gap**
  - ▶ Immediate send ACK, provided that segment starts at lower end of gap

# TCP Round Trip Time and Timeout

**Q:** how to set TCP timeout value?

- Longer than RTT
  - ▶ But RTT varies
- Too short: premature timeout
  - ▶ Unnecessary retransmissions
- Too long: slow reaction to segment loss

**Q:** how to estimate RTT?

- SampleRTT: measured time from segment transmission until ACK receipt
  - ▶ Ignore retransmissions
- SampleRTT will vary, want estimated RTT “smoother”
  - ▶ Average several recent measurements, not just current SampleRTT

# TCP Round Trip Time and Timeout

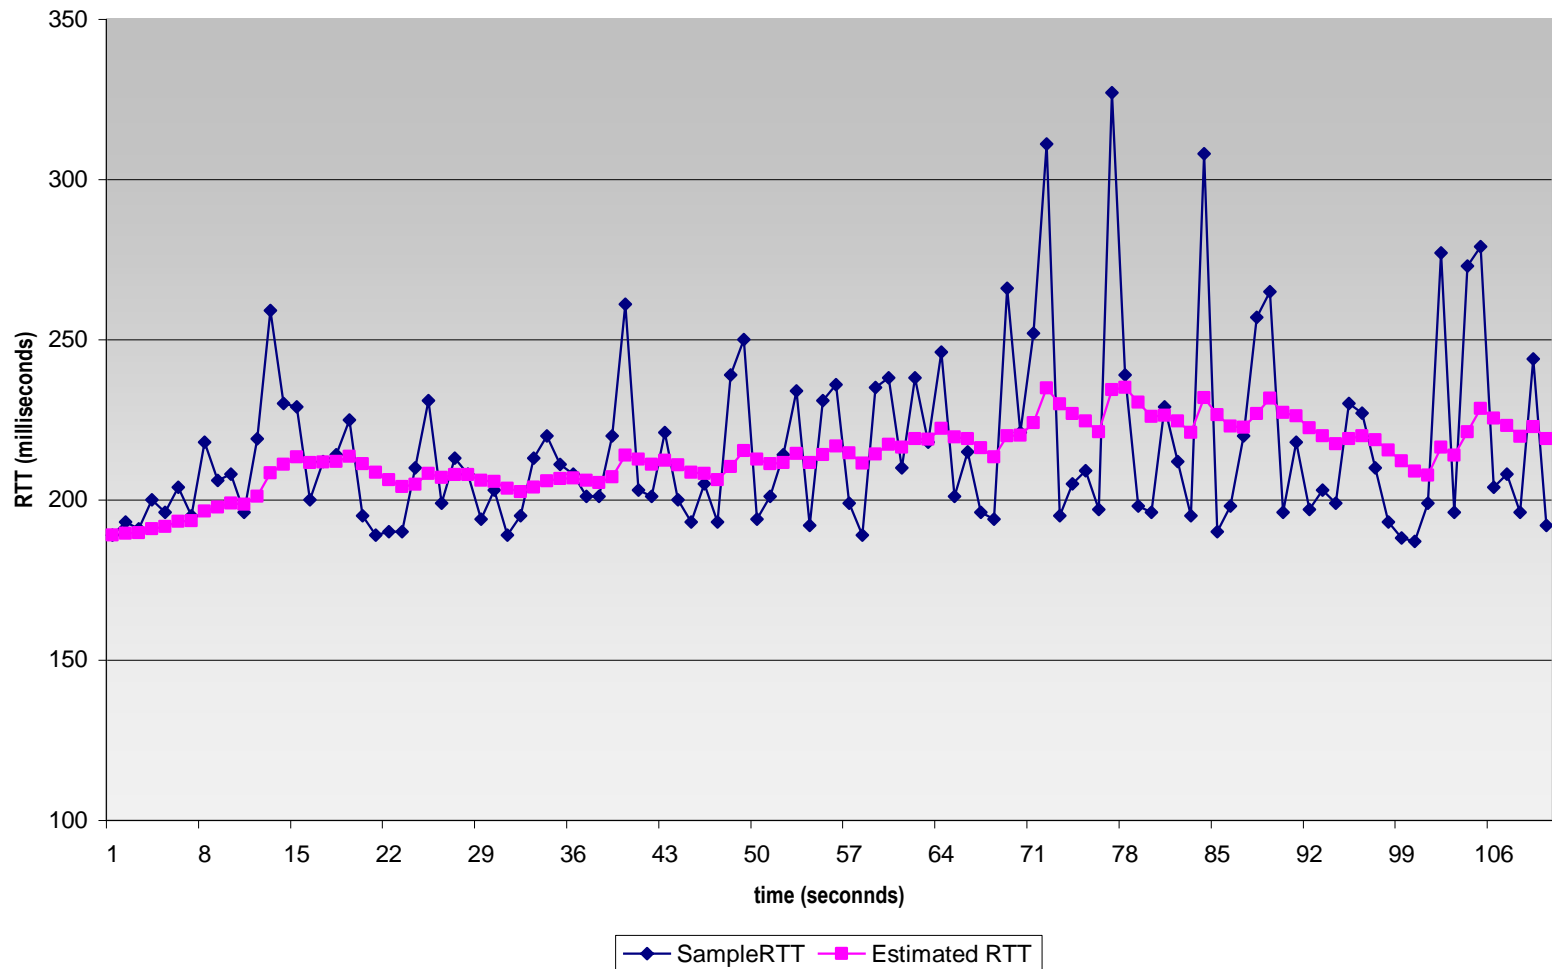
$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❑ Exponential weighted moving average
- ❑ Influence of past sample decreases exponentially fast
- ❑ Typical value:  $\alpha = 0.125$

# Example RTT Etimation:



RTT: gaia.cs.umass.edu to fantasia.eurecom.fr





# TCP Round Trip Time and Timeout

## Setting the timeout

- **EstimatedRTT plus “safety margin”**
  - ▶ Large variation in **EstimatedRTT** -> larger safety margin
- **Variance of RTT measurement:**
  - first estimate of how much **SampleRTT** deviates from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

### Then set timeout interval:

Avoid unnecessary retransmissions  
fast respond to changes

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

# TCP Connection Management



## Recall:

TCP sender, receiver establish “connection” before exchanging data segments

### ■ Initialize TCP variables:

- ▶ Seq. #s
- ▶ Buffers, flow control info (e.g. **RcvWindow**)

## Three way handshake:

**Step 1:** Client host sends TCP SYN segment to server

- ▶ Specifies initial seq #
- ▶ No data

**Step 2:** Server host receives SYN, replies with SYN ACK segment

- ▶ Server allocates buffers
- ▶ Specifies server initial seq. #

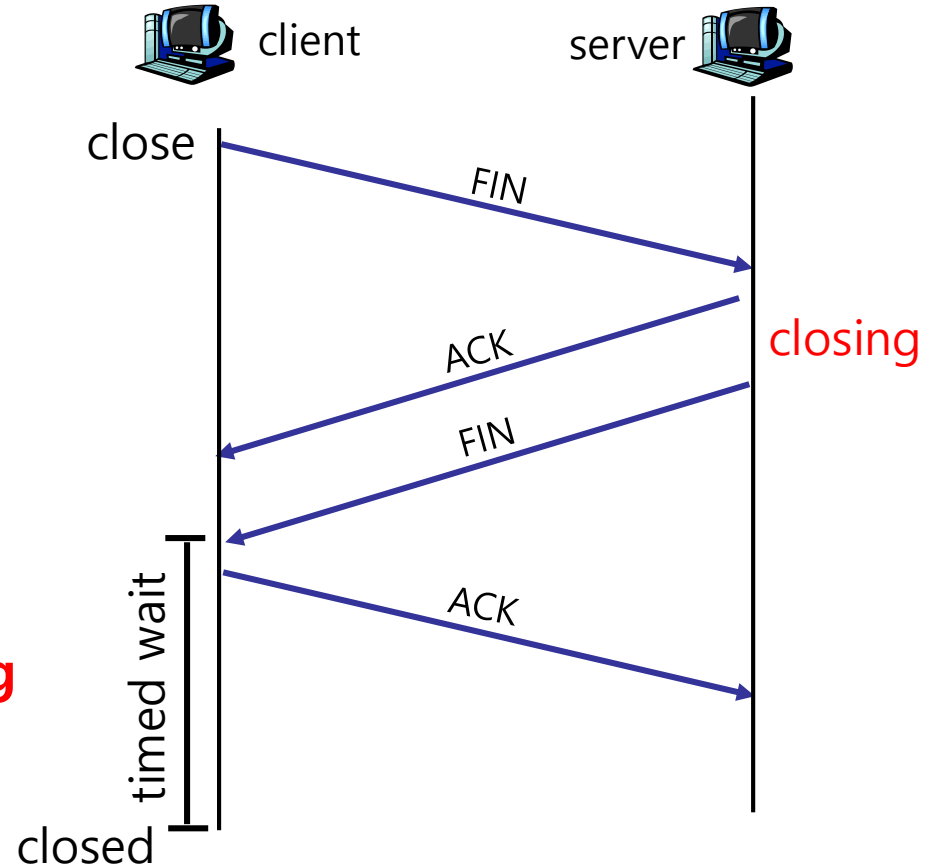
**Step 3:** Client receives SYN ACK, replies with ACK segment, which may contain data

# TCP Connection Management

## Closing a connection:

**Step 1:** Client end system sends TCP FIN control segment to server

**Step 2:** Server receives FIN, replies with ACK. **Inform Application. Send remaining data.** sends FIN.



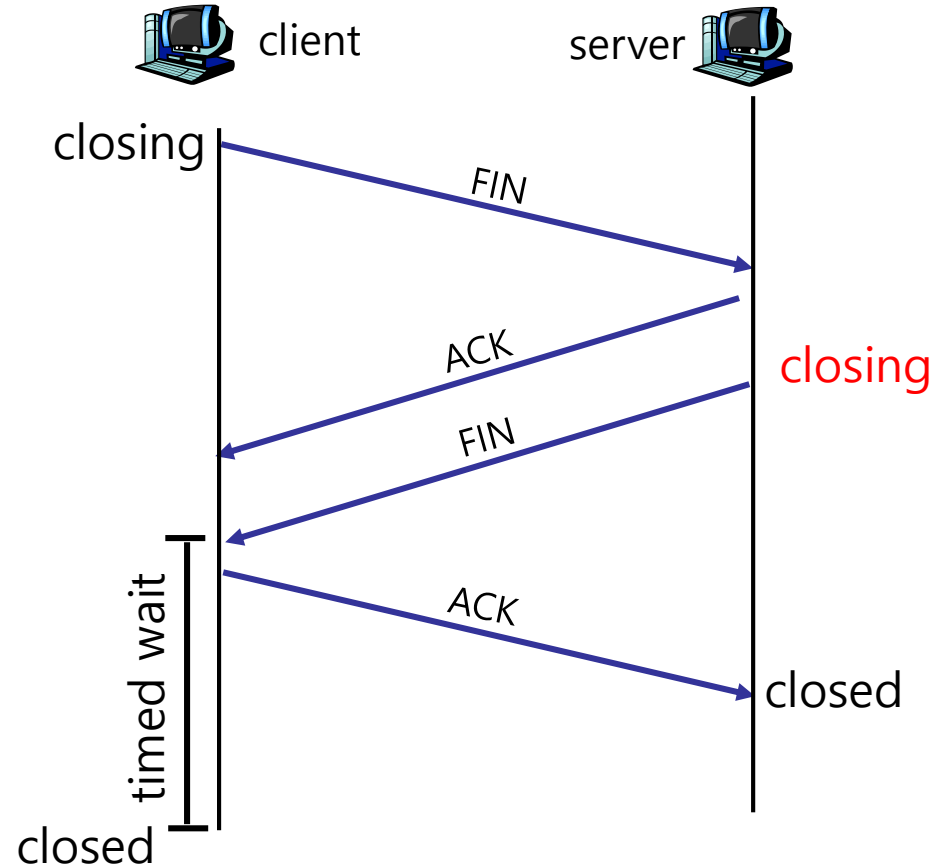
# TCP Connection Management

**Step 3:** Client receives FIN, replies with ACK.

- ▶ Enters "timed wait" - will respond with ACK to received FINs

**Step 4:** Server, receives ACK. Connection closed.

**Note:** with small modification, can handle simultaneous FINs.



# TCP Sender Events:

## Data received from app:

- Create segment with seq #
- Seq # is byte-stream number of first data byte in segment
- Start timer if not already running (think of timer as for oldest unacked segment)
- Expiration interval: `TimeoutInterval`

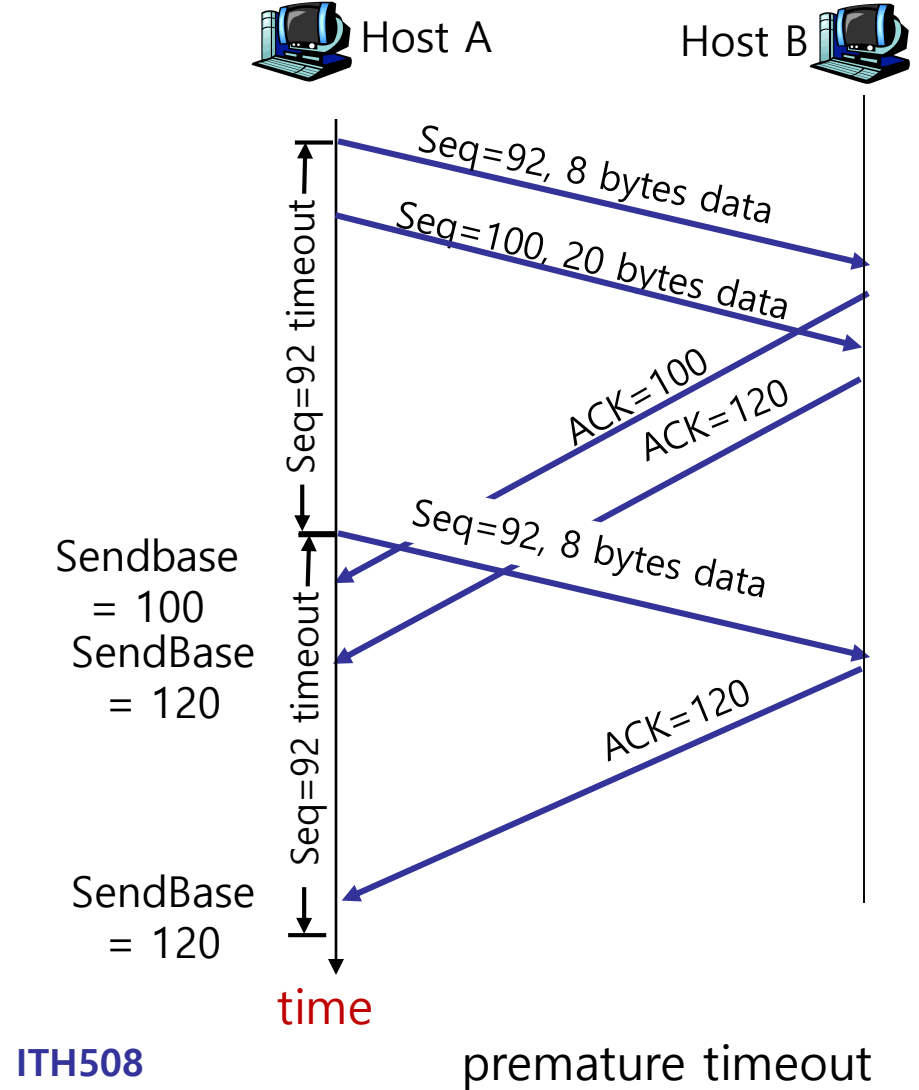
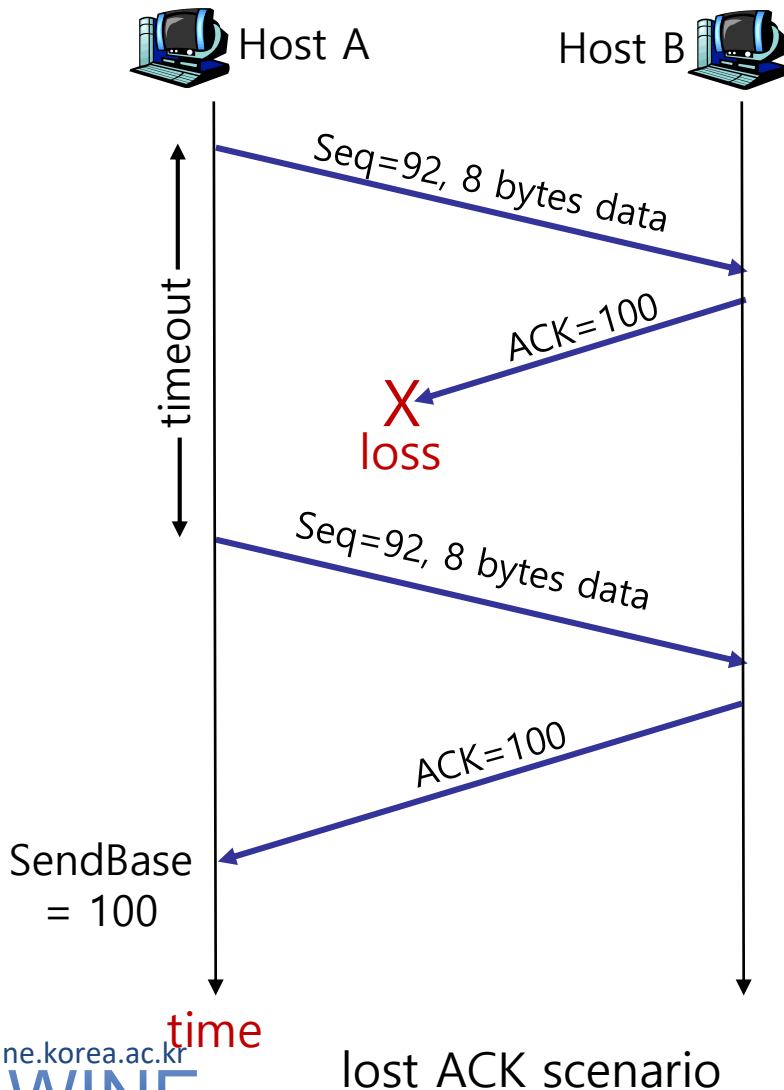
## Timeout:

- Retransmit segment that caused timeout
- Restart timer

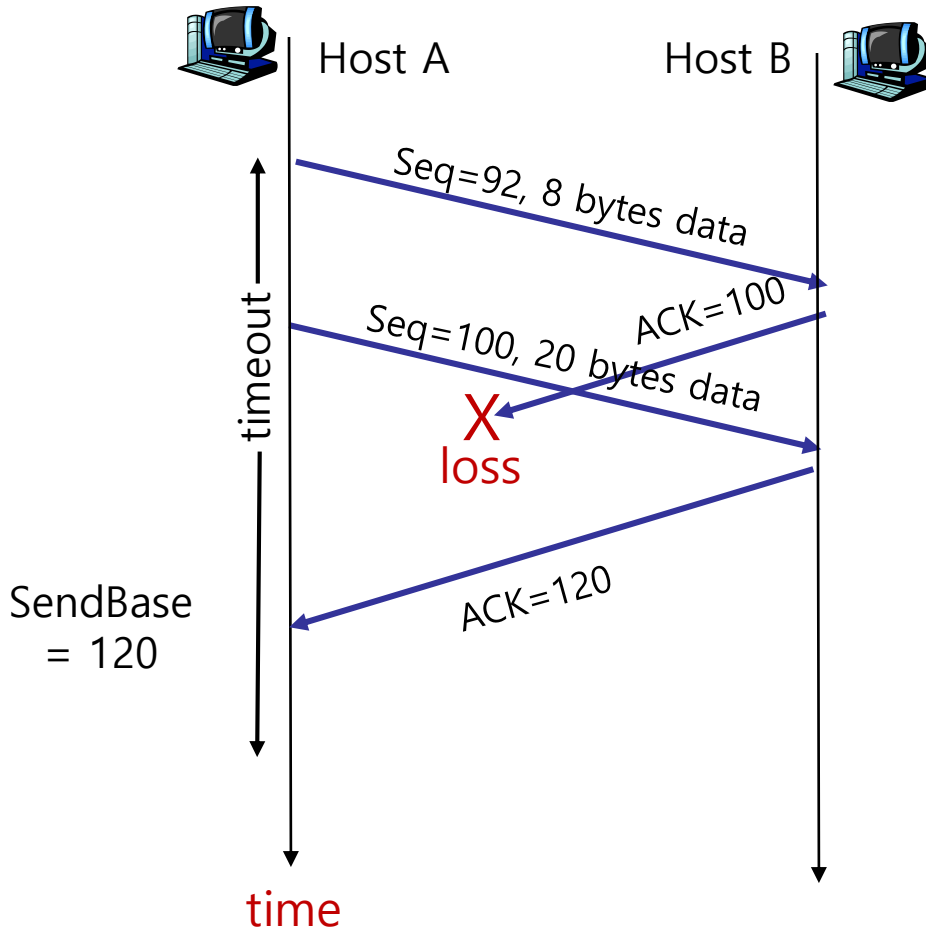
## Ack received:

- If acknowledges previously unacked segments
  - ▶ Update what is known to be acked
  - ▶ Start timer if there are outstanding segments

# TCP: Retransmission Scenarios



# TCP Retransmission Scenarios (more)



Cumulative ACK scenario

# TCP Flow Control vs. TCP Congestion Control



- **Flow control**

- ▶ Preventing senders from overrunning the capacity of the receivers

- **Congestion control**

- ▶ Preventing too much data from being injected into the network, **causing switches or links to become overloaded**

- **TCP provides both**

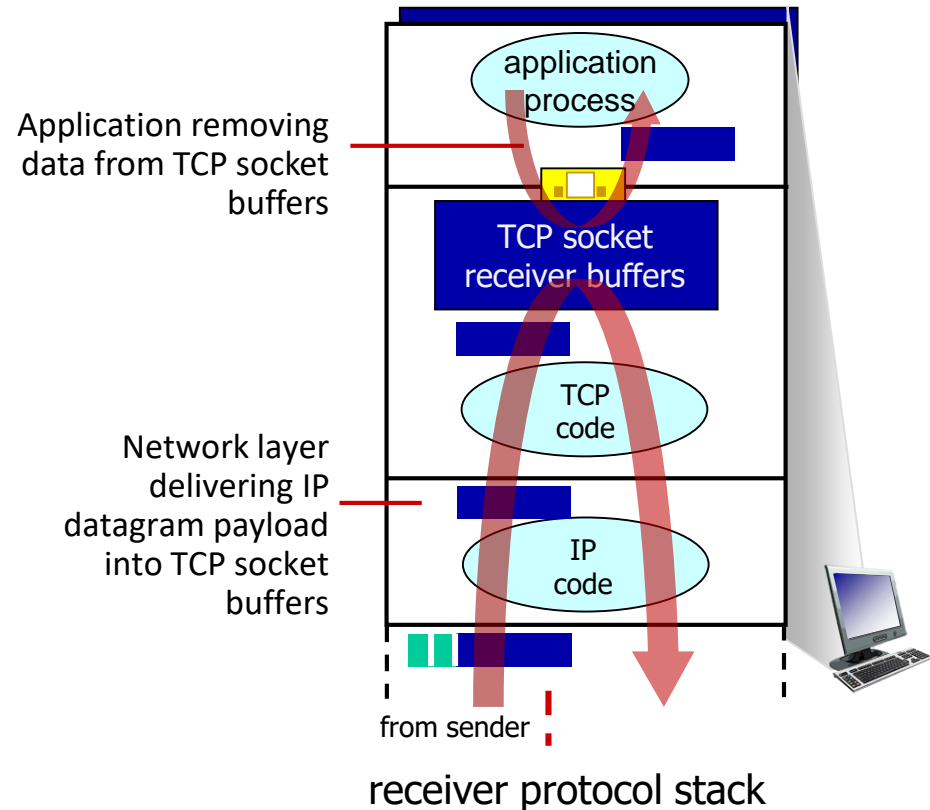
- ▶ Flow control based on advertised window
- ▶ Congestion control



# TCP Flow Control



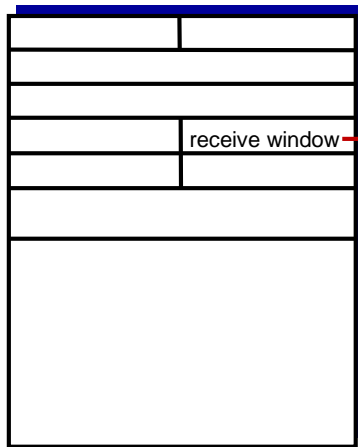
Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



# TCP Flow Control

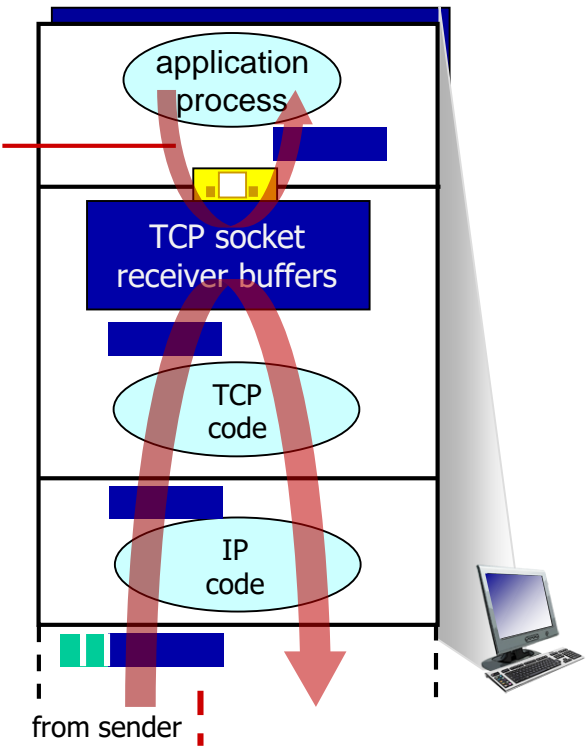


Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



flow control: # bytes receiver willing to accept

Application removing data from TCP socket buffers

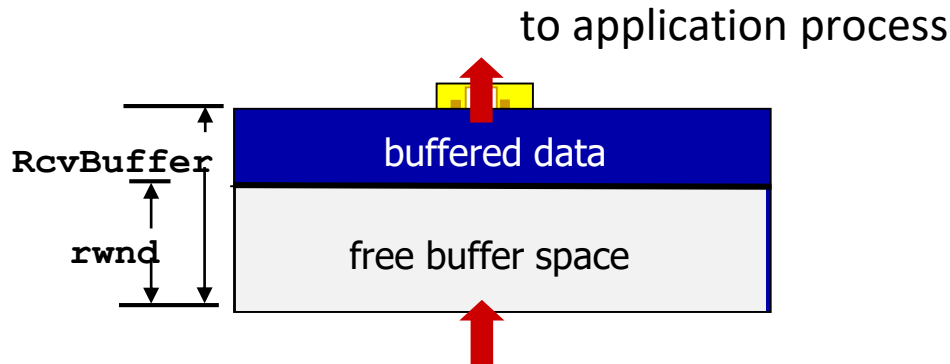


receiver protocol stack

# TCP Flow Control



- Receive side of TCP connection has a receive buffer:



TCP segment payloads

TCP receiver-side buffering

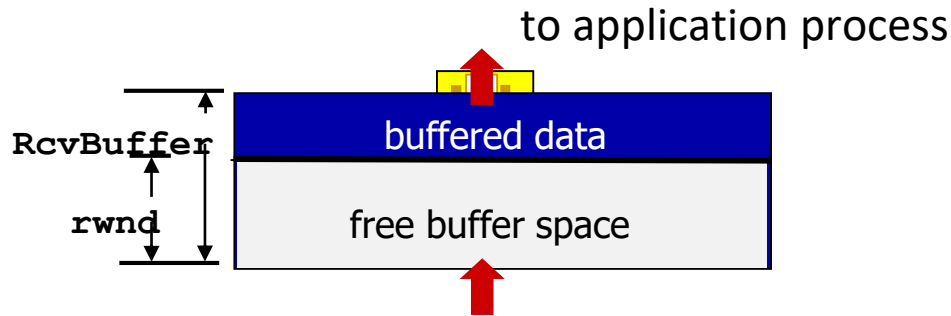
- Application process may be slow at reading from buffer

## flow control

Sender won't overflow receiver's buffer by transmitting too much, too fast

- Speed-matching service: matching the send rate to the receiving app's drain rate

# TCP Flow Control



TCP segment payloads

TCP receiver-side buffering

**(Suppose TCP receiver discards out-of-order segments)**

- **Spare room in buffer**
  - =  $\text{RcvWindow}$
  - =  $\text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$

- **Receiver advertises spare room by including value of RcvWindow in segments**
- **Sender limits unACKed data to RcvWindow**
  - ▶ Guarantees receive buffer doesn't overflow

# Congestion Control

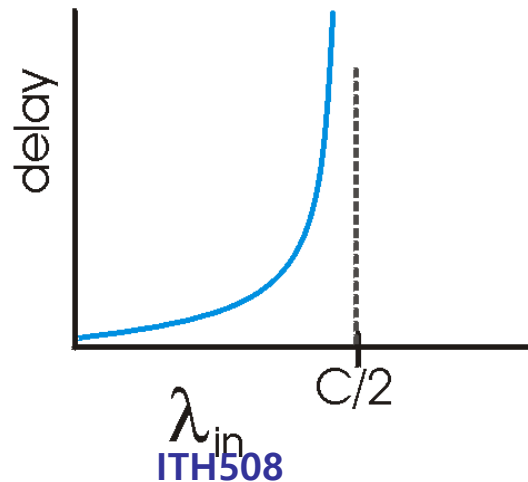
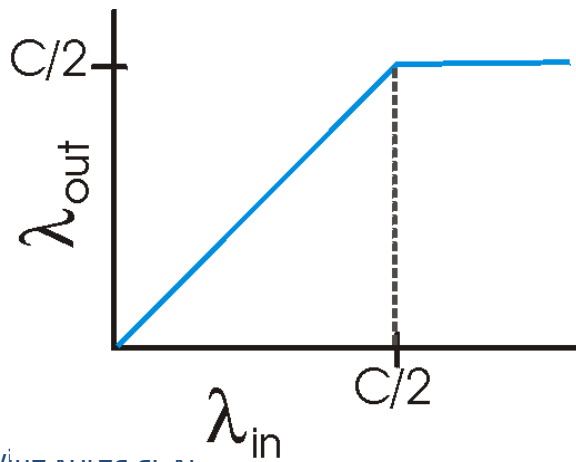
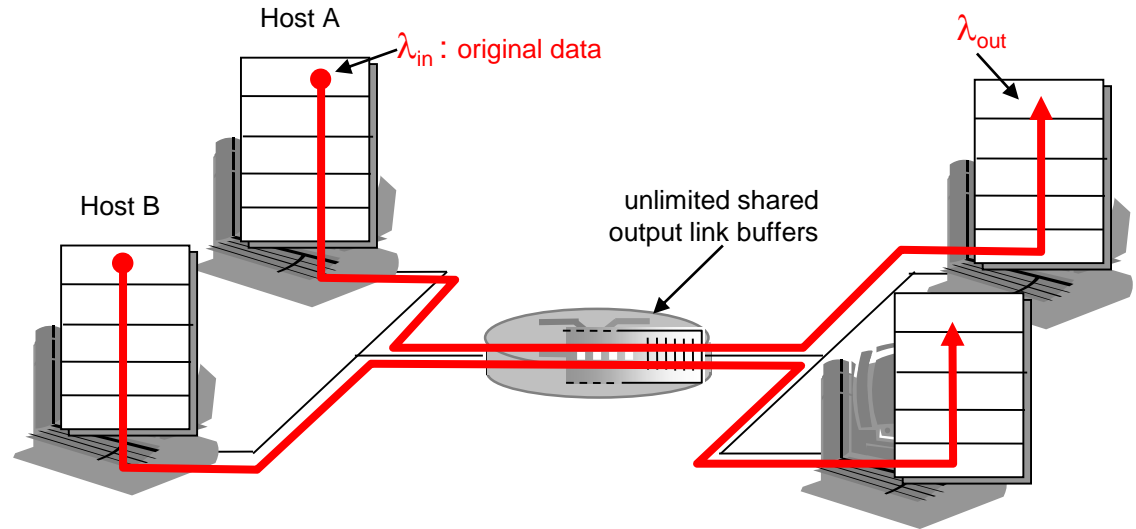


## Congestion:

- Informally: “Too many sources sending too much data too fast for *network* to handle”
- Different from flow control!
- Manifestations:
  - ▶ Lost packets (buffer overflow at routers)
  - ▶ Long delays (queueing in router buffers)

# Causes/Costs of Congestion: Scenario 1

- Two senders, two receivers
- One router, infinite buffers
- No retransmission

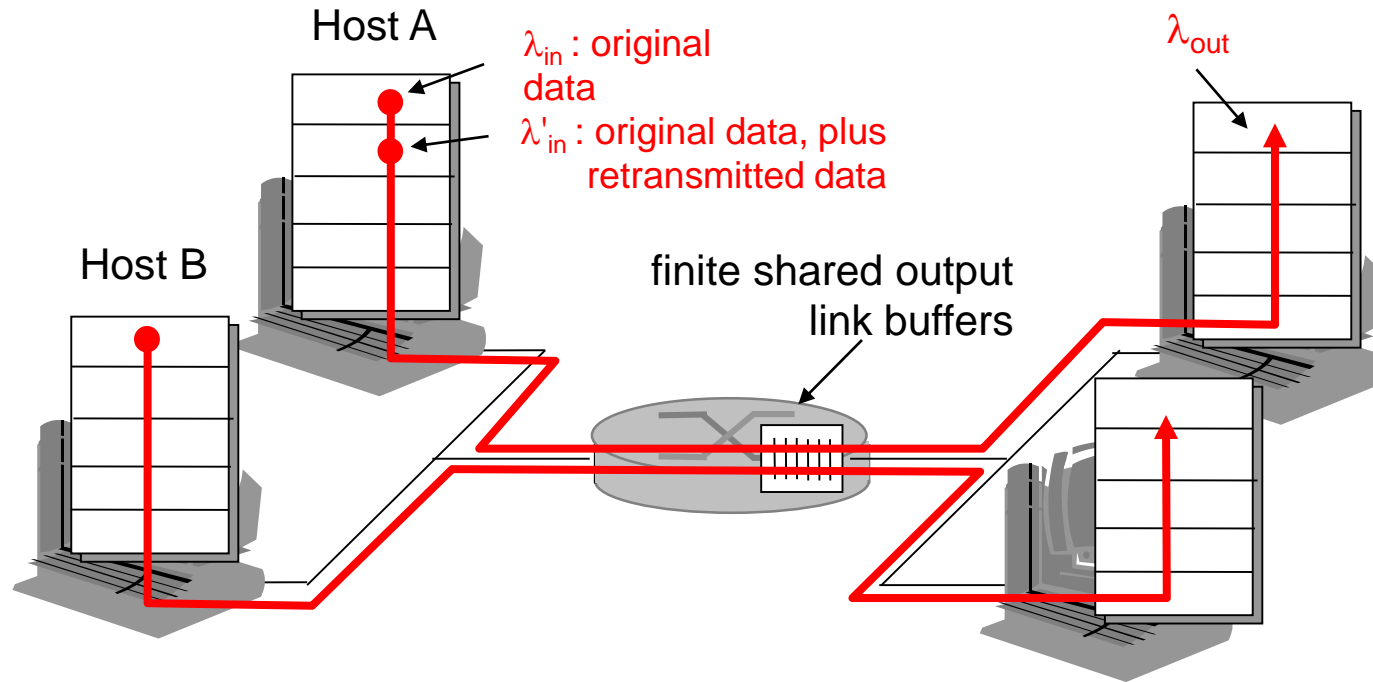


- Large delays when congested
- Maximum achievable throughput

# Causes/Costs of Congestion: Scenario 2



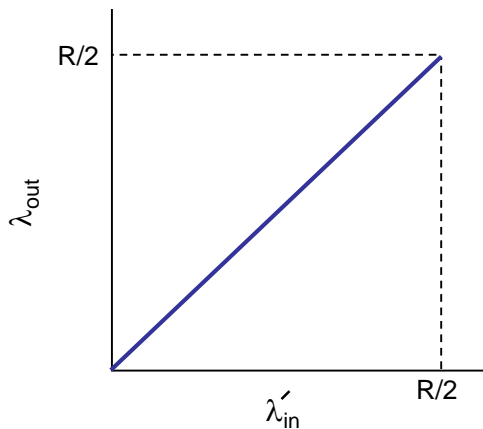
- One router, *finite* buffers
- Sender retransmission of lost packet



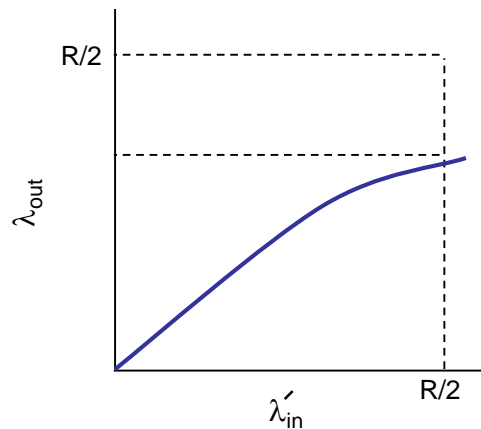
# Causes/Costs of Congestion: Scenario 2



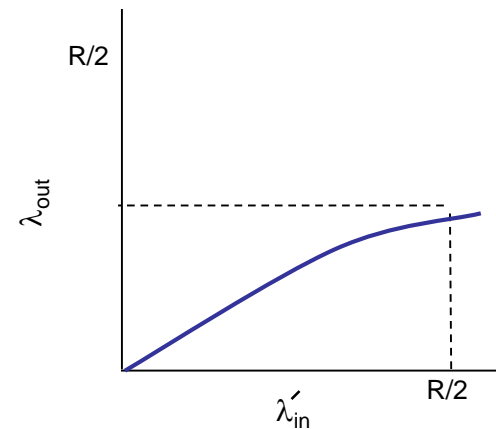
- a. Perfect case: always:  $\lambda_{in} = \lambda_{out}$  (goodput)
- b. Perfect retransmission only when loss:  $\lambda'_{in} > \lambda_{out}$
- c. Retransmission of delayed (not lost) packet makes  $\lambda'_{in}$  larger (than perfect case) for same  $\lambda_{out}$



a.



b.



c.

"costs" of congestion:

- ☐ More work (retransmission) for given "goodput"
- ☐ Unneeded retransmissions: link carries multiple copies of packet

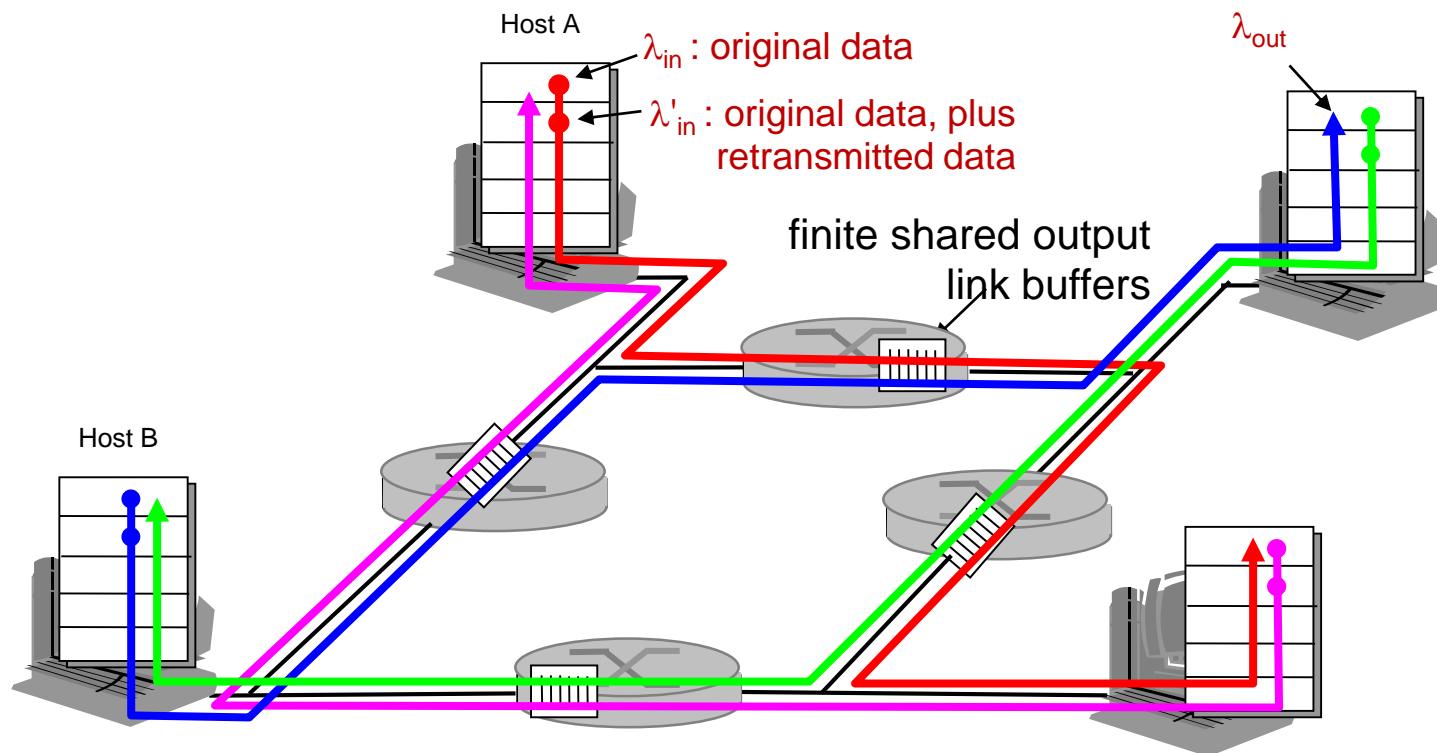


# Causes/Costs of Congestion: Scenario 3

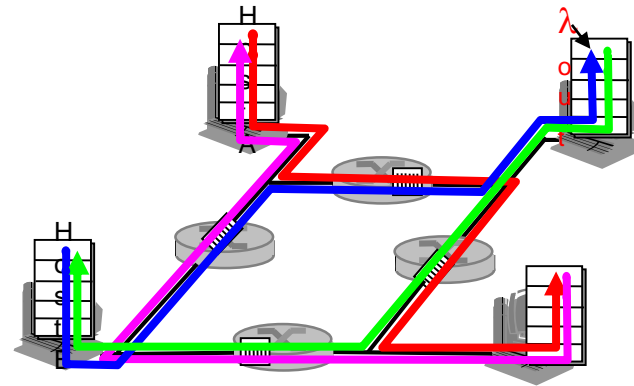
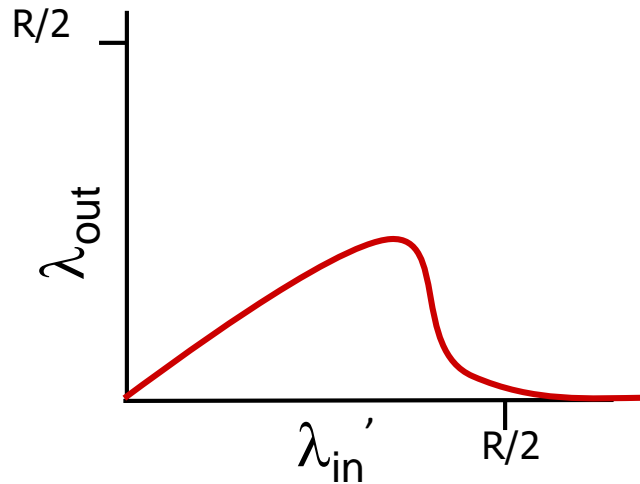


- Four senders
- Multihop paths
- Timeout/Retransmit

Q: what happens as  $\lambda_{in}$  and  $\lambda'_{in}$  increase ?



# Causes/Costs of Congestion: Scenario 3



Another "cost" of congestion:

- When packet dropped, any upstream transmission capacity used for that packet was wasted!

# TCP Congestion Control



## ■ Basic idea

- ▶ Add notion of congestion window
- ▶ Effective window (for transmission) **is smaller of**
  - **Advertised window** (flow control)
  - **Congestion window** (congestion control)
- ▶ Changes in congestion window size
  - **Slow increases** to absorb new bandwidth
  - **Quick decreases** to eliminate congestion

# TCP Congestion Control



## ■ Specific strategy

### ▶ Self-clocking

- Send data only when outstanding data is ACK'd
- Equivalent to send window limitation mentioned

### ▶ Initial time of TCP connection

- **Slow Start**
- Send data twice when outstanding data is ACK'd

# TCP Congestion Control



## ■ Specific strategy (continued)

### ▶ Growth

- Add one maximum segment size (MSS) per congestion window of data ACK'd
- Known as additive increase

### ▶ Decrease

- Cut window in half **when three duplicate ACKs**
- In practice, set window = window / 2
- Known as multiplicative decrease

### ▶ **Additive increase, multiplicative decrease (AIMD)**

- Congestion avoidance + Fast recovery

# TCP Start Up Behavior



## ■ How should TCP start sending data?

- ▶ AIMD is good for channels operating at capacity
- ▶ AIMD **can take a long time to ramp up to full capacity** from scratch
- ▶ Use Slow Start to increase window rapidly from a cold start

# TCP Start Up Behavior



- **Initialization of the congestion window**
  - ▶ Congestion window should start small
  - ▶ Avoid congestion due to new connections
  - ▶ **Start at 1 MSS, reset to 1 MSS with each timeout**
  - ▶ **Known as slow start**

# Slow Start



## ■ Objective

- ▶ Determine (probe) initial available capacity

## ■ Idea

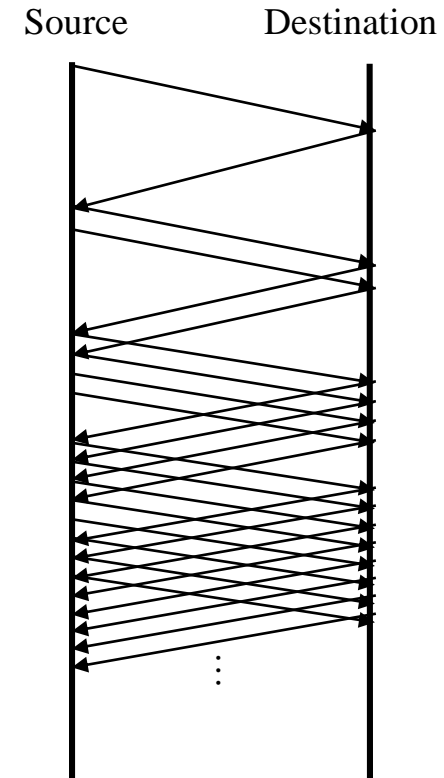
- ▶ Begin with **CongestionWindow** = 1 packet
- ▶ Double **CongestionWindow** each RTT
  - Increment by 1 packet for each ACK
- ▶ Continue increasing until loss

## ■ Result

- ▶ Exponential growth
- ▶ Slower than all at once

## ■ Used

- ▶ When first starting connection
- ▶ When connection times out





# Congestion Control



## ■ Question

- ▶ How **does the source determine whether or not the network is congested?**

## ■ Answer

- ▶ **Timeout** signals packet loss
  - **Serious situation**
- ▶ **3 duplicate ACK** means also packet loss
  - **Not that serious congestion situation**
- ▶ **Packet loss is rarely due to transmission error (on wired lines)**
- ▶ **Lost packet implies congestion!**

# Congestion Control

## ■ Congestion indication

- ▶ Time out
- ▶ 3 duplicate ACKs

## ■ How to respond to losses

- ▶ Reset and slow-start
- ▶ Fast Retransmit
- ▶ Fast Recovery
- ▶ Selective ACK

## ■ Result

- ▶ Congestion avoidance + Slow Start

# Fast Retransmit



- Time-out period often relatively long:
  - ▶ Long delay before resending lost packet
- Detect lost segments via duplicate ACKs.
  - ▶ Sender often sends many segments back-to-back
  - ▶ If segment is lost, there will likely be many duplicate ACKs.
- If sender receives 3 ACKs for the same data, it supposes that segment after ACKed data was lost:
  - ▶ **Fast retransmit**: resend segment before timer expires

# Fast Recovery



## ■ Fast recovery

- ▶ When **fast** retransmission occurs, **skip slow start**
- ▶ Congestion window becomes  $1/2$  previous
- ▶ Start additive increase immediately

# Additive Increase/Multiplicative Decrease



## ■ Objective

- ▶ Adjust to changes in **available capacity** (not maximal)

## ■ Tools

- ▶ React to observance of congestion (time out, 3 duplicate ACKs)
- ▶ Probe channel to detect more resources

## ■ Observation

- ▶ On notice of congestion
  - Decreasing too slowly will not be reactive enough
  - Sometimes, need to reset TCP congestion dynamics
- ▶ On probe of network
  - Increasing too quickly will overshoot limits

# Additive Increase/Multiplicative Decrease



## ■ New TCP state variable

### ▶ CongestionWindow

- Similar to **AdvertisedWindow** for flow control

### ▶ Limits how much data source can have in transit

- $\text{MaxWin} = \text{MIN}(\text{CongestionWindow}, \text{AdvertisedWindow})$
- $\text{EffWin} = \text{MaxWin} - (\text{LastByteSent} - \text{LastByteAcked})$
- TCP can send no faster than the slowest component, network or destination

## ■ Idea

- ▶ Increase **CongestionWindow** when congestion goes down
- ▶ Decrease **CongestionWindow** when congestion goes up

# Additive Increase/Multiplicative Decrease



## ■ Algorithm

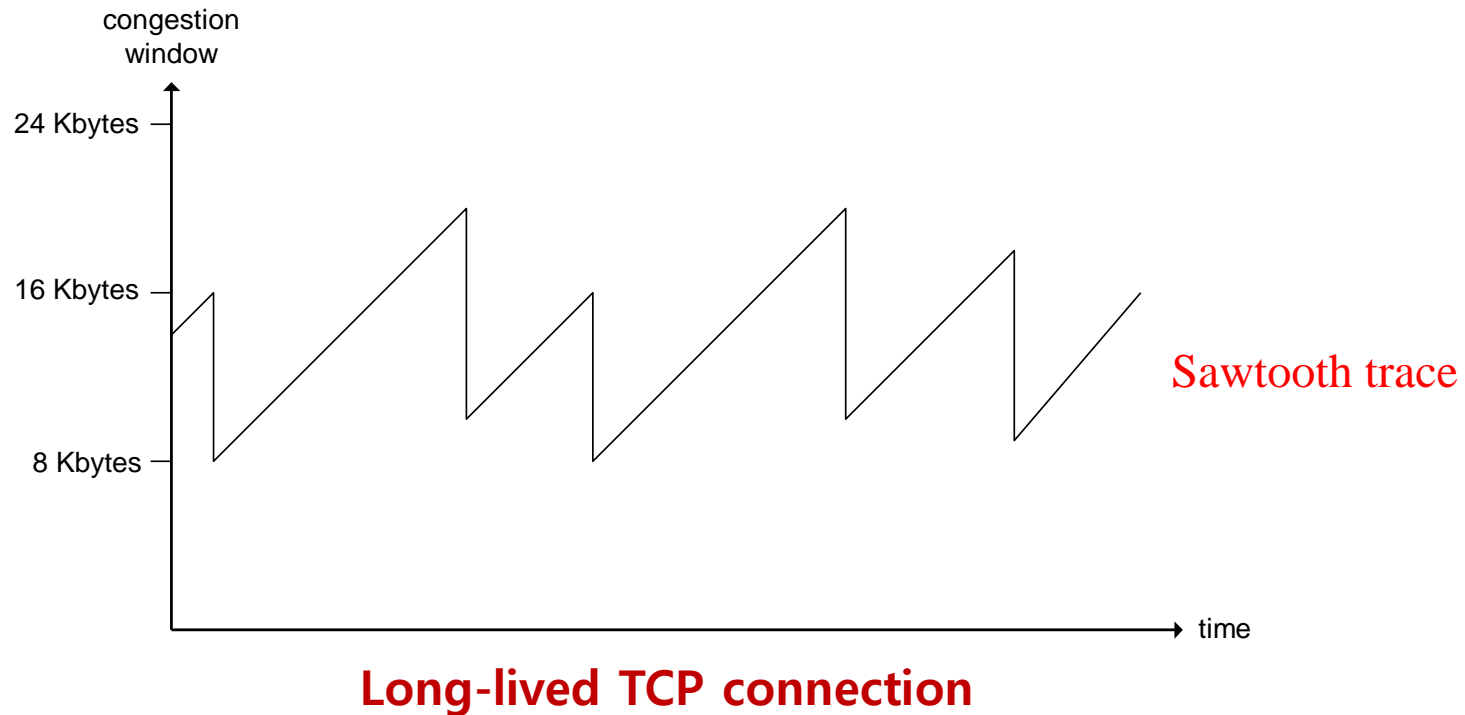
- ▶ Increment CongestionWindow by **one packet** per RTT
  - Linear increase
- ▶ Divide CongestionWindow by two whenever a congestion occurs
  - Multiplicative decrease

# Additive Increase/Multiplicative Decrease



**Multiplicative decrease:**  
cut CongWin in half  
after loss event

**Additive increase:** increase  
CongWin by 1 MSS every  
RTT in the absence of loss  
events: *probing*





# Interaction between Slow Start and AIMD



- Maintain **threshold window size**
- Use multiplicative increase
  - ▶ When congestion window smaller than threshold
  - ▶ **Double window** for each window ACK'd
- Threshold value
  - ▶ Initially set **to maximum window size**
  - ▶ **Set to 1/2 of current window on timeout**
- In practice, increase congestion window by one MSS for each ACK of new data (or N bytes for N bytes)

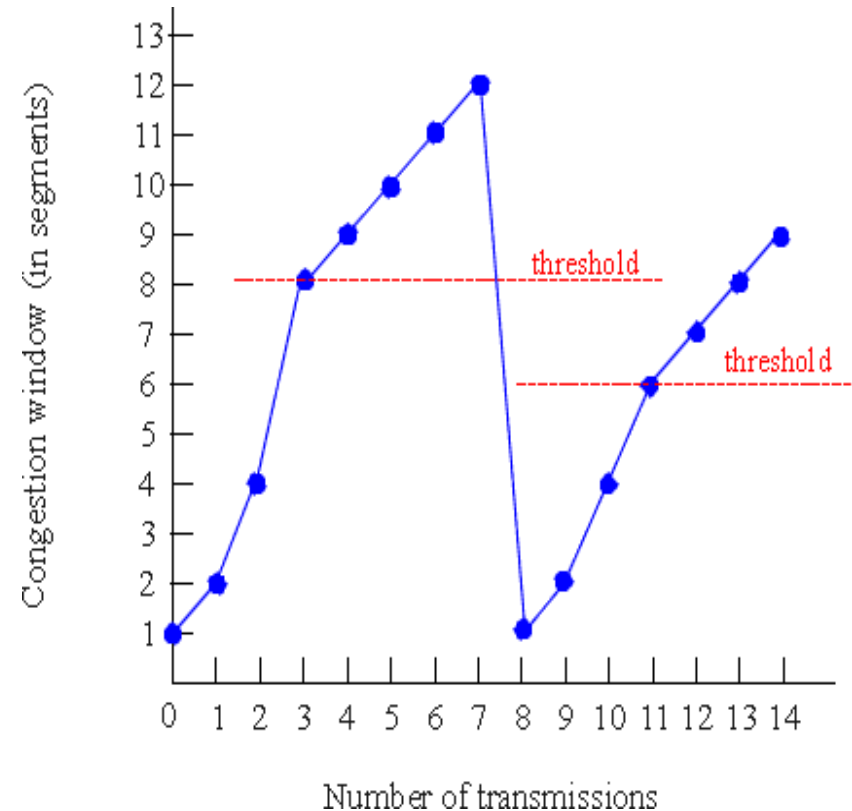
# Interaction between Slow Start and AIMD

## ■ How to control the exponential increase?

- ▶ Target window size is **ssthresh**
- ▶ Estimate network capacity
- ▶ When **CongestionWindow** reaches **ssthresh**
  - Switches to additive increase

## ■ Example

- ▶ Initial values
  - **ssthresh** = 8
  - CongestionWindow = 1
- ▶ Loss after transmission 7
  - CongestionWindow currently 12
  - Set **ssthresh** = CongestionWindow/2
  - Set CongestionWindow = 1



# TCP Congestion Control Mechanism

- End-end control (no network assistance)

- Sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{CongWin}$$

- Roughly,

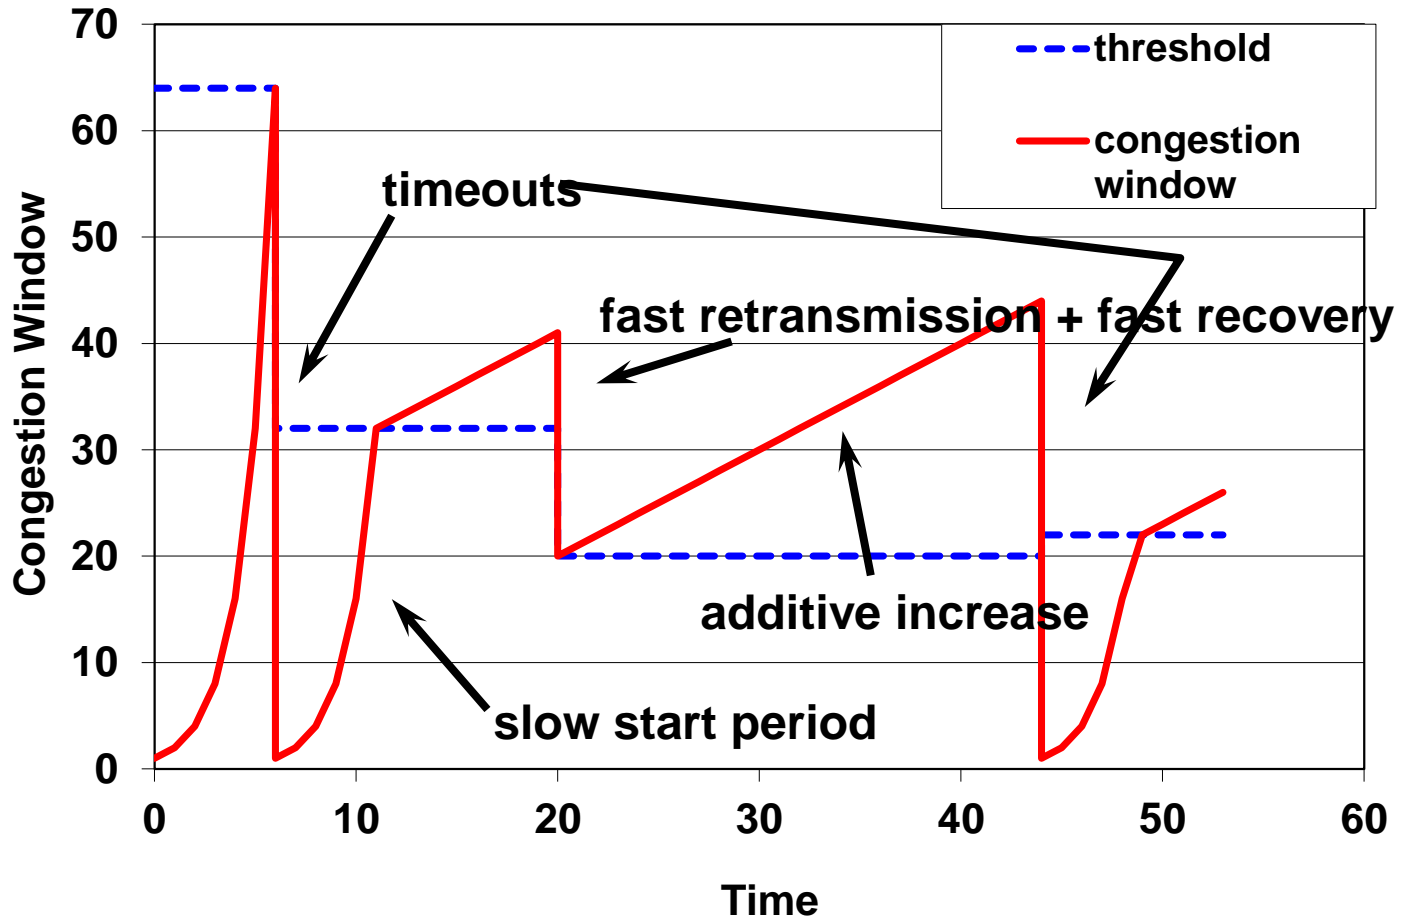
$$\text{rate} = \frac{\text{CongWin}}{\text{RTT}} \text{ Bytes/sec}$$

- CongWin is dynamic, function of perceived network congestion

## Main mechanisms:

- ▶ Slow start
- ▶ AIMD
- ▶ Conservative after timeout events
- ▶ Congestion window (cwnd)-based
- ▶ Congestion indication

# TCP Congestion Window Trace



# TCP Variants



## ■ TCP Tahoe

- ▶ Use 3 duplicate ACKs as well as Timeout
- ▶ Fast Retransmit

## ■ TCP Reno

- ▶ Fast Retransmit
- ▶ Fast Recovery

## ■ TCP New Reno

- ▶ Fast Retransmit
- ▶ Fast Recovery
- ▶ Partially ACK

## ■ TCP SACK

- ▶ SACK permitted
- ▶ TCP reno

## ■ TCP Cubic

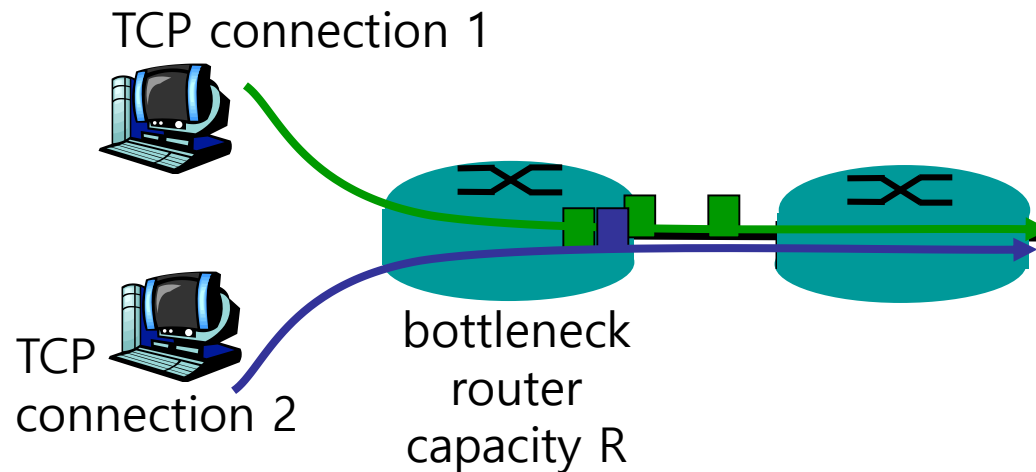
- ▶ It increases  $W$  as a function of the cube of the distance between the current time and when TCP window size will reach the maximum window
  - larger increases when further away from the maximum window
  - smaller increases (cautious) when nearer the maximum window

# TCP FAIRNESS

# TCP Fairness



**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$



# TCP Fairness



## Two competing sessions:

- Additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally

